

[← Back to CCAR-F Tests](#)

Agentic Architecture & Orchestration - Test Complete!

Here's how you performed on this CCAR-F domain

90%

Overall Score

9

Correct Answers

9:38

Time Taken

Excellent! You're well-prepared for this domain.

This was a sample test from a limited pool of questions.
Upgrade to Premium for access to the complete CCAR-F question
bank with 250+ questions!

Upgrade to Premium

Get unlimited access to:

- ✔ Unlimited practice tests across all 5 domains
- ✔ Full-length mock exams
- ✔ AI-powered explanations and study guides
- ✔ Progress tracking and performance analytics
- ✔ Smart practice focusing on weak areas

Question Review

Review your answers and explanations

Show Explanations

- 1 A support agent calls several account tools before composing a customer response. The current application stops whenever Claude's text contains the word "completed," causing it to terminate before some requested tools run. Which implementation BEST improves reliability? ✔

A. Continue until two consecutive responses contain no tool names.

B. Execute requests on `tool_use`, append matching tool results, and continue until `end_turn`.
Correct

C. Begin a new conversation after each tool call to reduce context size.

D. Require Claude to end every final response with a unique completion phrase.

- 2 A research agent can search regulatory filings, query market data, and retrieve internal reports. The correct next source depends on evidence returned by the previous tool, but production policy requires a maximum of eight tool calls. Which orchestration approach is MOST appropriate? ✔

A. Always call the three tools in a fixed order and stop after one pass.

B. Let Claude invoke tools without retaining results or enforcing a call limit.

C. Let Claude select each next tool while the host preserves results and enforces the limit.
Correct

D. Ask Claude to list every tool it might need before the host executes any.

3 An investment research system has specialists for patents, financial statements, and industry news. Investigations vary by company, findings can conflict, and the final report must synthesize evidence across sources. Which architecture BEST meets these requirements?



A. Use a coordinator to select specialists, aggregate results, and request targeted refinement.

Correct

B. Have every specialist exchange unrestricted messages with every other specialist until consensus.

C. Run every specialist once and concatenate their responses in a predefined section order.

D. Assign the entire investigation to one specialist selected from the company industry code.

4 A repository-wide authentication change affects an API handler, a shared token library, integration tests, and deployment documentation. The handler depends on the revised library interface, while the documentation and initial test analysis can proceed independently. Which decomposition strategy is BEST?



A. Give each subagent several isolated code lines, then merge all edits by timestamp.

B. Let all subagents modify every affected file concurrently and resolve conflicts automatically.

C. Complete all code, tests, and documentation serially with one general-purpose subagent.

D. Assign coherent concerns, sequence interface-dependent work, parallelize independent work, then run integration review. Correct

5 A Claude-based coding coordinator has an AgentDefinition for a dependency-audit subagent, but the coordinator cannot launch it. The



coordinator's allowedTools contains Read and Grep but not Task. The subagent also needs the target package path and approved-license list. What should the architect change?

- A. Add Bash to the coordinator and let the subagent discover all requirements.
- B. Add Task to allowedTools and pass the package path and license constraints explicitly.
Correct
- C. Move the subagent instructions into the coordinator's system prompt and remove its definition.
- D. Grant every available tool to both agents and rely on repository documentation for context.

- 6 A support agent may issue refunds above \$1,500 only after identity verification and policy approval. Skipping either step creates material financial risk. If approval is unavailable, a human reviewer must receive a complete handoff. Which design BEST satisfies these requirements? 

- A. Tell Claude to verify and approve first, then trust its generated explanation as evidence.
- B. Allow the refund immediately, but send the transaction to a reviewer for later confirmation.
- C. Enforce both checks before execution and hand off completed checks, facts, reason, and unresolved issues. **Correct**
- D. Ask the customer to confirm twice, then treat the repeated confirmation as policy approval.

- 7 A support agent uses a payment tool whose amount field may be returned as "\$42.10," "4210 cents," or a numeric decimal. Refund calls must also be blocked when the verified account ID differs from the requested account ID. Which implementation is MOST appropriate? 

- A. Add examples to the prompt and ask Claude to check account IDs and normalize amounts.
Your Answer

B. Normalize only in the reporting database and audit account mismatches at the end of each day.

C. Let each downstream subagent interpret the amount and independently decide whether to refund.

D. Intercept mismatched calls before execution and normalize successful results in PostToolUse.
Correct

8 A production defect initially appears in `invoice_writer.py`, but early inspection may reveal dependencies in schema adapters, retry logic, or database migrations. The full investigation scope is unknown, and the final diagnosis must explain cross-file interactions. Which strategy is BEST?



A. Use a fixed chain that always analyzes the same four files in the same order.

B. Analyze relevant files locally, adapt the plan as dependencies emerge, then perform cross-file synthesis. **Correct**

C. Send the entire repository to several agents and accept the first diagnosis returned.

D. Analyze `invoice_writer.py` only because expanding scope could increase latency and token use.

9 A named Claude Code session was interrupted after analyzing six files for a checkout refactor. Overnight, another developer substantially changed two analyzed files, while the other four remained unchanged. The team wants to preserve useful context without treating stale file observations as current. What should the architect do FIRST?



A. Resume the named session and re-analyze the two changed files before continuing.
Correct

B. Resume unchanged and instruct Claude to trust every result already in the session.

C. Fork the session solely to make all earlier tool results current automatically.

D. Delete all session state and restart without any summary of the prior investigation.

10 A customer-support agent can call tools to retrieve an order, check delivery status, and issue a replacement. A single request may require several tool calls. The application currently ends the loop whenever the assistant text contains the phrase "case resolved," but this sometimes terminates processing before a requested tool is executed. Which implementation BEST controls the loop? 

A. Continue until the assistant produces a response longer than a configured minimum length.

B. Execute requests on tool_use, append each tool result, and continue until end_turn.
Correct

C. Require the assistant to place the phrase "case resolved" in a dedicated XML element.

D. Execute every available support tool once, then request a final summary from the assistant.



Take Test Again

Try Other Domains