



Anthropic

CCAR-P

Claude Certified Architect - Professional

Version: Demo

[Total Questions: 10]

Web: www.dumpspedia.com

Email: support@dumpspedia.com

IMPORTANT NOTICE

Feedback

We have developed quality product and state-of-art service to ensure our customers interest. If you have any suggestions, please feel free to contact us at feedback@dumpsmedia.com

Support

If you have any questions about our product, please provide the following items:

- 🔍 exam code
- 🔍 screenshot of the question
- 🔍 login id/email

please contact us at support@dumpsmedia.com and our technical experts will provide support within 24 hours.

Copyright

The product of each order has its own encryption code, so you should use it independently. Any unauthorized changes will inflict legal punishment. We reserve the right of final explanation for this statement.

Category Breakdown

Category	Number of Questions
Stakeholder Communication and Lifecycle Management	3
Evaluation, Testing and Optimization	2
Integration	1
Solution Design and Architecture	1
Claude Models, Prompting and Context Engineering	1
Developer Productivity and Operational Enablement	2
TOTAL	10

Question #:1 - [Stakeholder Communication and Lifecycle Management]

You are reviewing a peer's draft architecture decision record (ADR). The ADR states the decision and the technical detail but does not state the alternatives considered or the trade-offs accepted.

Which response is most appropriate?

- A. Delete the ADR entirely and rely on the engineering team's collective memory to reconstruct the decision context, alternatives, and trade-offs during future audits or onboarding.
- B. Ask the author to add the alternatives considered and the trade-offs accepted before publishing the ADR.
- C. Approve the ADR as written on the grounds that the decision itself is documented, accepting that reviewers and future maintainers will lack the context to evaluate or revisit it.
- D. Replace the ADR with a different decision unrelated to the one drafted.

Answer: B

Explanation

An ADR must preserve not only the selected decision but also the context, alternatives evaluated, and trade-offs consciously accepted. Without this information, future engineers cannot determine whether the decision remains valid after requirements, Claude models, costs, security controls, or integration capabilities change. Option B completes the record while retaining the valid technical detail already documented. Option A destroys useful evidence and replaces durable documentation with unreliable organizational memory. Option C produces an incomplete record that cannot support audits, onboarding, or later architectural review. Option D changes the decision without evidence or a decision-making process. For Claude solutions, ADRs should typically capture the chosen model and integration pattern, rejected alternatives, evaluation evidence, security implications, operational consequences, and conditions that would trigger reconsideration.

=====

Question #:2 - [Evaluation, Testing and Optimization]

A Claude-based research assistant begins producing responses that confidently contradict its retrieved source documents despite no change to the retrieval pipeline.

Which two diagnostic actions most directly identify the root cause of this behavior? (Select two.)

- A. Increase the context-window size to allow more retrieved chunks per query.
- B. Switch the retrieval index to a denser embedding model to improve chunk-relevance scores.
- C. Determine whether the failure reproduces on the previous model version to test for a model mismatch.
- D. Reduce the temperature setting to lower response variance across all query types.
- E. Inspect the system-prompt grounding instructions to determine whether citation constraints remain intact.

Answer: C E

Explanation

The retrieval pipeline is stated to be unchanged, so the first investigation should isolate changes in generation behavior and grounding instructions. Replaying the failing cases on the previous model version determines whether the regression follows the model upgrade or current model-prompt combination. Inspecting the system prompt establishes whether source-use rules, uncertainty behavior, citation requirements, or prohibitions against unsupported claims were removed or weakened. Increasing context or replacing the embedding model changes a component for which no failure evidence exists and may introduce additional noise. Lowering temperature may reduce variation but does not explain systematic contradiction of retrieved evidence. The investigation should compare identical requests, retrieved passages, prompts, and model versions before applying a fix. Define success criteria and evaluations

Question #:3 - [Integration]

A technical team is debating whether to implement a new capability for generating regulatory filings from internal data as a Claude Skill or as an MCP server. The capability requires a fixed authoring procedure with embedded examples, does not need to call live external systems, and must be portable across three Claude deployments: claude.ai, an internal API integration, and Claude Code.

Which factor most strongly favors implementing the capability as a Skill?

- A. The capability needs to query a live database for the most recent regulatory filings.
- B. The capability must also be callable from a developer's command-line interface.
- C. The capability requires direct authentication against the corporate identity provider.
- D. The capability is procedural knowledge with no live external-system calls required.

Answer: D

Explanation

A Skill is the appropriate packaging mechanism for reusable procedural knowledge, instructions, templates, examples, and supporting resources. The filing capability follows a defined authoring method and requires no live external-system interaction, so an MCP server would introduce unnecessary transport, authentication, hosting, and operational complexity. MCP becomes appropriate when Claude must invoke live databases, APIs, or enterprise services. Command-line availability does not independently determine the design because Skills can be used across supported Claude environments. Direct identity-provider authentication would favor an integration or MCP service rather than a self-contained Skill. Anthropic describes Agent Skills as portable packages that provide agents with specialized procedural and organizational expertise. Agent Skills

Question #:4 - [Evaluation, Testing and Optimization]

A Claude Architect is reviewing a post-deployment performance report for an AI-assisted legal-document summarization system. The report includes these observations:

Average summarization time decreased from 47 minutes to 6 minutes per document.

Associates spend less time on summaries, but overall billable output has not measurably changed.

Infrastructure costs increased by 22% because redundant retry logic generated additional API calls.

Some summaries require attorney correction, adding an average of 8 minutes of review per document.

Which analysis correctly attributes each observation to the appropriate business-value pillar?

- A. Observations 1 and 2 both indicate efficiency gains; Observation 3 is a solution-cost issue; Observation 4 is a performance-SLA issue.
- B. Observation 1 is a transformation outcome; Observation 2 is an efficiency gain; Observation 3 is a performance-SLA degradation; Observation 4 is a solution-cost issue.
- C. Observations 1 and 4 together indicate a net performance-SLA improvement; Observation 2 is a transformation gap; Observation 3 is a productivity drain from over-engineering.
- D. Observation 1 indicates an efficiency gain; Observation 2 shows that productivity has not yet been realized; Observation 3 is a solution-cost issue; Observation 4 is an efficiency loss that partially offsets Observation 1.

Answer: D

Explanation

Observation 1 is a direct efficiency improvement because the time required for one summarization task fell substantially. Observation 2 shows that local time savings have not yet translated into higher organizational productivity or billable output. Observation 3 is explicitly a solution-cost problem caused by unnecessary API calls. Observation 4 introduces additional human correction time and therefore partially offsets the gross efficiency gain reported in Observation 1. Option D is the only answer that keeps these value dimensions distinct. Measuring only generation speed would overstate realized value because review effort, retry costs, and downstream output must also be included. A valid post-deployment assessment should calculate net cycle time, cost per accepted summary, correction rate, and resulting business output.

Question #:5 - [Stakeholder Communication and Lifecycle Management]

A senior architect is preparing briefing materials on a new retrieval architecture. The executive sponsor has requested a summary of the architectural decision. Which framing is most appropriate for that audience?

- A. Technical alternatives evaluated, implementation implications, and component-level consequences.
- B. Business outcomes achieved, trade-offs accepted, and high-level risks acknowledged.
- C. Capabilities delivered, scope implications, and feature-level dependencies on the roadmap.
- D. Threat model, control mappings, residual-risk acceptance, and audit traceability.

Answer: B

Explanation

Executive sponsors generally need architecture decisions translated into business significance. The briefing should explain what outcomes the retrieval architecture enables, why the chosen approach is valuable, which material trade-offs were accepted, and what major risks require awareness or sponsorship. This allows the executive to assess strategic alignment, investment value, and organizational exposure without becoming immersed in component-level implementation details. Option A is more appropriate for architects and engineering teams conducting a technical review. Option C is primarily suited to product and program stakeholders managing features, scope, and roadmap dependencies. Option D is designed for security, compliance, and risk-governance audiences requiring detailed control evidence. Although those perspectives remain important, option B provides the concise, outcome-oriented framing needed for an executive architectural decision summary.

=====

Question #:6 - [Solution Design and Architecture]

A healthcare organization is evaluating two Claude-powered AI architectures for a clinical documentation assistant. Architecture X produces higher output quality scores but costs \$0.18 per documentation session and averages 4.2 seconds per response. Architecture Y scores slightly lower on quality metrics but costs \$0.09 per session and averages 2.1 seconds per response. The stated SLA requires responses under 3 seconds, and the annual volume is projected at 2 million documentation sessions.

Which evaluation approach correctly applies business value pillar analysis to this decision?

- A. Select Architecture Y based solely on the 50% cost reduction, since solution cost is the most important value pillar in healthcare budget-constrained environments.
- B. Select Architecture X because the higher quality scores justify the cost premium, and any SLA gap can be addressed through infrastructure optimization after deployment.

- C. Eliminate Architecture X on SLA grounds, then evaluate Architecture Y against the efficiency and solution cost pillars by calculating annual cost difference and assessing whether the quality delta materially impacts clinical workflow productivity.
- D. Recommend a hybrid approach using Architecture X for complex cases and Architecture Y for routine cases, without additional analysis, since this preserves quality where it matters most.

Answer: C

Explanation

Architecture X fails a stated production gate: its 4.2-second average exceeds the required three-second response SLA. It should therefore not remain a qualifying architecture unless it is redesigned and retested before selection. A promise to optimize after deployment is not evidence of compliance.

Architecture Y meets the latency requirement at 2.1 seconds. Its annual inference cost is \$180,000 at two million sessions, compared with \$360,000 for Architecture X—a difference of \$180,000 annually. However, cost alone does not establish suitability. The organization must determine whether Y's lower quality score causes material omissions, correction effort, clinician dissatisfaction, documentation risk, or reduced workflow productivity.

Option C correctly treats the SLA as a threshold and then analyzes the remaining architecture across cost, efficiency, quality, and productivity. Option A ignores clinical quality. Option B disregards an explicit SLA failure. Option D may eventually be viable, but it introduces routing complexity and leaves unanswered how complex cases are detected, whether X's latency becomes acceptable for them, and how both paths are evaluated.

The final choice should be gated by clinical-quality evaluation, human review requirements, safety metrics, operational latency, and total annual cost.

Study Guide references/topics: Business-value pillars; performance gates; solution cost; workflow efficiency; quality trade-offs; healthcare architecture evaluation.

=====

Question #:7 - [Claude Models, Prompting and Context Engineering]

You are reviewing a peer's draft system prompt that contains contradictory instructions: one section says never to speculate beyond the supplied source, while another says to confidently fill in any gaps.

Which response is most appropriate?

- A. Add a priority instruction directing the model to evaluate all instructions and apply whichever appears most contextually appropriate on each request.
- B. Remove or rewrite the gap-filling instruction so the prompt consistently constrains the model to source-supported content.

- C. Increase temperature so output randomness masks the contradiction.
- D. Keep both instructions and rely on the model to decide which one to follow on each request.

Answer: B

Explanation

The prompt contains mutually incompatible behavioral requirements. The architect should resolve the contradiction explicitly so the model receives one coherent grounding policy. If the application requires source-supported answers, the prompt should prohibit unsupported completion, permit an insufficient-evidence response, and require citations or supporting passages. Adding another priority instruction leaves the underlying conflict in place and creates ambiguity about when speculation is permitted. Increasing temperature makes behavior more variable rather than more correct. Retaining both instructions produces inconsistent outputs across requests and model versions. Anthropic recommends clear, direct, specific prompting and notes that conflicting project instructions can cause Claude to select among them unpredictably. Prompting best practices; project-instruction troubleshooting

Question #:8 - [Stakeholder Communication and Lifecycle Management]

You are running a discovery session with a business sponsor who asks for “an AI system to fix our customer escalations.”

Which approach best yields actionable requirements?

- A. Accept the sponsor’s original phrasing as a complete and final requirement and proceed directly to design without asking targeted questions to surface actors, triggers, or constraints.
- B. Decompose the request by asking targeted questions about who escalates, what triggers escalation, what good resolution looks like, and which constraints—latency, cost, audit, and data sensitivity—apply.
- C. Defer the discovery session indefinitely on the grounds that the sponsor must fully specify the system before the architect can ask any clarifying or scoping questions.
- D. Assume the requirement matches the architect’s previous engagement in a different industry and proceed with that context, without validating actors, triggers, or constraints for this sponsor.

Answer: B

Explanation

The sponsor’s statement expresses a desired outcome but does not define an implementable requirement. Option B converts the vague request into concrete information about users, workflow triggers, expected resolutions, decision boundaries, and non-functional constraints. These findings can then be translated into measurable success criteria, evaluation cases, integration requirements, and escalation rules. Anthropic’s evaluation guidance states that an LLM application should begin with specific, measurable, achievable, and relevant success criteria. Option A prematurely treats an aspiration as a complete specification. Option C incorrectly transfers discovery responsibility entirely to the sponsor, while Option D imports assumptions

from an unrelated context. Effective Claude architecture requires traceability from the business problem through requirements, evaluation measures, design decisions, and production acceptance criteria. Anthropic: Define success criteria and build evaluations

=====

Question #:9 - [Developer Productivity and Operational Enablement]

You are identifying inefficiency in a Claude Code workflow where each engineer manually re-explains the project 's conventions and architecture in every session.

Which adjustment most directly removes this inefficiency?

- A. Forbid the use of Claude Code on the project entirely to avoid the session-initialization overhead, accepting that the team loses all AI-assisted development productivity for this codebase.
- B. Capture the project 's conventions and architecture in a project-scoped CLAUDE.md (or equivalent persistent project-context file) committed to the repository, so each session loads it automatically.
- C. Tell each engineer to retype the project conventions and architecture context more quickly at the start of each session, reducing time lost without eliminating the repeated manual effort.
- D. Remove all documented project conventions and architectural standards so engineers have nothing to re-explain at session start, accepting the loss of consistency and shared coding standards.

Answer: B

Explanation

A project-scoped CLAUDE.md converts repeatedly supplied information into persistent, version-controlled context. Claude Code loads applicable CLAUDE.md instructions at session initialization, allowing architecture principles, coding conventions, build commands, repository structure, and verification requirements to be shared consistently across engineers. Anthropic explicitly recommends using CLAUDE.md for information that developers would otherwise explain repeatedly; concise and specific instructions provide stronger adherence. Committing the file also enables normal review and change-control practices. Option C reduces typing speed but preserves the underlying waste and inconsistency. Options A and D remove useful capability or governance rather than correcting the workflow. The implementation should periodically review the file for obsolete, conflicting, or unnecessarily verbose instructions.

Study Guide references/topics: Claude Code project memory; CLAUDE.md; persistent context; repository conventions; workflow standardization.

Question #:10 - [Developer Productivity and Operational Enablement]

You are preparing an operational runbook for a Claude-based service.

Which content is essential to include in the runbook?

- A. Dashboard and log references only, without alert definitions, triage steps, escalation paths, or rollback procedures for the on-call engineer to act on.
- B. Common alerts and their triage steps, escalation paths, rollback procedures, and references to the relevant dashboards and logs.
- C. Alert definitions and triage steps only, without escalation paths, rollback procedures, or references to dashboards and logs for on-call use.
- D. Escalation paths and rollback procedures only, without alert definitions, triage steps, or dashboard references to guide initial incident response.

Answer: B

Explanation

An actionable runbook must connect detection, diagnosis, containment, recovery, and escalation. Option B supplies that complete operational chain: alerts identify abnormal conditions; triage procedures guide initial investigation; dashboards and logs provide evidence; escalation paths assign authority; and rollback procedures restore a known-safe state. The runbook should cover model errors, latency degradation, rate limits, tool failures, prompt regressions, cost anomalies, data-governance events, and dependency outages. Anthropic's operational materials demonstrate incident workflows in which alerts lead responders to logs and runbooks before corrective action. Options A, C, and D each omit information required for an on-call engineer to progress from notification to controlled resolution. Procedures should be tested through exercises and updated after incidents.

Study Guide references/topics: Claude operational and observability patterns; incident response; triage; escalation; rollback; operational readiness.

About dumpspedia.com

dumpspedia.com was founded in 2007. We provide latest & high quality IT / Business Certification Training Exam Questions, Study Guides, Practice Tests.

We help you pass any IT / Business Certification Exams with 100% Pass Guaranteed or Full Refund. Especially Cisco, CompTIA, Citrix, EMC, HP, Oracle, VMware, Juniper, Check Point, LPI, Nortel, EXIN and so on.

View list of all certification exams: [All vendors](#)

Microsoft



CITRIX



EMC²
where information lives[®]


(ISC)²



JUNIPER
NETWORKS



ORACLE



vmware

We prepare state-of-the art practice tests for certification exams. You can reach us at any of the email addresses listed below.

- Sales: sales@dumpspedia.com
- Feedback: feedback@dumpspedia.com
- Support: support@dumpspedia.com

Any problems about IT certification or our products, You can write us back and we will get back to you within 24 hours.