


**Isha Training Solutions****CCAR-P • CLAUDE CERTIFIED ARCHITECT
PROFESSIONAL MOCK EXAMS**[Overview](#)[Mock
Exams](#)[+91-9133190573](#)[Start a mock exam](#)

Independent, unofficial exam prep by Isha Training Solutions – not affiliated with,
endorsed by, or sponsored by Anthropic.

[Home](#) / [AI Courses](#) / [CCAR-P: Claude Certified Architect Professional Mock Exams](#)**1000**/**1000**

21 of 21 correct on Mock Exam 2

[Retake this exam](#)[Back to mock exams](#)[Review your answers](#)

✔ Correct • Integration

A RAG system handles natural-language questions well but performs poorly on queries containing exact part numbers and error codes. What retrieval improvement addresses this?

- ☒ A Adding hybrid search that combines keyword/exact-match retrieval with semantic search, so exact-term queries are still matched precisely. Correct answer
- ☐ B Removing the retrieval step entirely.
- ☐ C Switching to pure semantic/vector search exclusively.
- ☐ D Increasing the chunk size to include the entire corpus in one chunk.

WHY

Pure semantic search can struggle with exact identifiers; combining it with keyword-based matching (hybrid search) covers both fuzzy conceptual queries and precise exact-term lookups.

✔ Correct • Integration

After an initial retrieval step returns 50 candidate passages, a reranking step is added before the top few are passed to the model. What does this improve?

- ☐ A It guarantees the retrieval index never needs updating again.
- ☐ B It has no effect on output quality.
- ☒ C It improves the precision of what actually reaches the model, by re-scoring the initial candidates more carefully so the truly most relevant ones are prioritized. Correct answer
- ☐ D It removes the need for an initial retrieval step altogether.

WHY

Reranking sits between broad initial retrieval and the model call, refining a wider candidate set down to a more precisely ordered short list — improving what actually ends up in the model's context.

✔ Correct • Integration

An integration currently authenticates using a broad, personal credential belonging to one engineer. What is the architectural concern, and the better pattern?

- ☒ **A** It conflates two very different risk profiles and complicates auditing/revocation; a dedicated, appropriately scoped service account is the safer, clearer pattern for automated integrations. Correct answer
- ☐ **B** There is no concern; personal credentials are equivalent to service accounts.
- ☐ **C** Personal credentials always have narrower scope than service accounts.
- ☐ **D** Service accounts cannot be used for automated systems.

WHY

Reusing a personal, broadly-scoped credential for an automated integration makes access harder to audit and revoke independently of the person, whereas a scoped service account cleanly separates that risk.

✔ Correct • Integration

Which of the following are sound practices for architecting a resilient integration against a third-party dependency that occasionally goes down? (Select all that apply.)

- ☒ **A** Designing a graceful degradation path (e.g., a fallback response or reduced functionality) rather than failing the whole user request. Correct answer
- ☐ **B** Retrying the dependency indefinitely with no timeout or limit whenever it's unavailable.
- ☒ **C** Instrumenting the integration with tracing so a future outage can be diagnosed quickly. Correct answer
- ☒ **D** Using a dedicated, appropriately scoped service account instead of a broad personal credential for the integration. Correct answer

WHY

Graceful degradation, properly scoped credentials, and tracing/observability are all sound resilience practices. Retrying indefinitely with no limit or timeout is itself a risk, not a resilience practice.

✔ Correct • Solution Design and Architecture

A process has well-defined, rarely-changing steps with clear rules for what happens next at each stage. Does this favor a fixed workflow or a dynamic, tool-calling agent architecture?

- ☐ A A dynamic agent architecture, since agents are always superior to workflows.
- ☐ B The choice should be made without regard to how well-defined the steps are.
- ☒ C **A fixed workflow, since well-defined and stable steps don't need the flexibility (and added complexity) of dynamic, model-driven decision-making at each step.** Correct answer
- ☐ D Neither approach is usable for well-defined processes.

WHY

When steps are well-defined and stable, a simpler fixed workflow is usually the better fit — the flexibility of a dynamic agent architecture earns its complexity when steps are genuinely uncertain or variable, not when they're already fixed.

✔ Correct • Solution Design and Architecture

A team is deciding whether to build a custom retrieval-evaluation pipeline in-house or adopt an existing, well-supported open-source one. What should primarily guide this build-vs-buy decision?

- ☐ A The decision should be made without evaluating existing options at all.
- ☐ B Always build custom, regardless of what already exists.
- ☐ C Always buy or adopt existing tools, regardless of fit.

- ☒ D **Whether the existing option genuinely meets the requirements and is well-supported — if so, adopting it is usually more efficient than duplicating that effort in-house.** Correct answer

WHY

A sound build-vs-buy call weighs whether an existing, well-supported solution actually fits the need — reflexively building custom everything (or reflexively adopting everything) skips the actual evaluation that should drive the decision.

✔ Correct · Solution Design and Architecture

A task naturally decomposes into several distinct, independently-improvable responsibilities (e.g., style review, security review, test coverage review). What architectural choice suits this better than one generalist agent trying to do all three?

- ☒ **A** A fleet of narrow, specialized subagents, each targeted and independently evaluable/improvable for its specific responsibility. Correct answer
- ☐ **B** Removing two of the three responsibilities to simplify the system.
- ☐ **C** Task decomposability has no bearing on this architectural choice.
- ☐ **D** A single generalist orchestrator handling all three responsibilities identically with no specialization.

WHY

When a task cleanly decomposes into distinct responsibilities, narrow specialized subagents are typically easier to target, evaluate, and improve independently than one generalist trying to do everything at once.

✔ Correct · Evaluation, Testing and Optimization

Before shipping a prompt change to a production feature, what should the evaluation dataset used to test it look like?

- ☐ **A** No evaluation dataset is necessary before shipping.
- ☐ **B** The exact same examples used to originally write the prompt.
- ☐ **C** A small, convenient set of examples the engineer happens to have on hand.
- ☒ **D** A held-out set that's genuinely representative of real production traffic, so results generalize to what users actually send. Correct answer

WHY

An evaluation set only tells you something useful if it actually represents real production traffic — a convenient but unrepresentative sample can hide regressions that show up only with real user input.

✔ **Correct** · Evaluation, Testing and Optimization

A team suspects hallucination rates have crept up after several small changes, but has only been checking output manually and occasionally. What's a more scalable way to detect this?

- ☐ A Only check outputs when a user complains.
- ☒ B **Add automated evaluation — such as a grading model or rule-based checks — that can systematically score a large sample of outputs for hallucination rate over time.** Correct answer
- ☐ C Continue relying entirely on occasional manual spot checks.
- ☐ D Stop checking altogether since hallucinations can't be measured.

WHY

Manual spot checks don't scale and can miss gradual drift. Automated evaluation — a grader model or rule-based check applied to a larger, consistent sample — gives a systematic, trackable signal instead.

✔ **Correct** · Evaluation, Testing and Optimization

A structured data-extraction task and an open-ended long-form writing task both need evaluation. Should they use the same evaluation metric?

- ☐ A No evaluation metric can be applied to either task type.
- ☐ B Yes, always use exact-match scoring for every task type.
- ☐ C Always use a single generic score for every task with no customization.

- ☒ **D** No — exact-match scoring suits the extraction task well, while the open-ended writing task is better served by a rubric-based or model-graded evaluation that can judge quality along multiple dimensions.

Correct answer

WHY

The right metric depends on the task: exact-match works well for tasks with one correct structured answer, while open-ended writing needs a richer, rubric- or grader-based approach that can judge quality, not just exact correctness.

✔ **Correct** • Evaluation, Testing and Optimization

Which of the following make an evaluation program for a production Claude feature genuinely trustworthy? (Select all that apply.)

- ☒ **A** Using a held-out evaluation set that's representative of real production traffic.
- ☒ **B** Choosing an evaluation metric that actually fits the task, such as a rubric-based grader for open-ended writing rather than exact-match scoring.
- ☒ **C** Continuously monitoring quality, latency, and cost after launch, not just at initial evaluation time.
- ☐ **D** Relying only on a handful of convenient examples the engineer already had on hand.

Correct answer

Correct answer

Correct answer

WHY

A representative held-out set, ongoing production monitoring, and a task-appropriate metric are what make evaluation genuinely trustworthy. A convenient but unrepresentative sample undermines exactly that.

✔ **Correct** • Governance, Safety and Risk Management

A production system makes automated decisions that occasionally affect customers materially. Why does the system need an audit trail for these decisions?

- ☒ **A** So any consequential decision can be traced and explained after the fact, supporting accountability and enabling review when something goes wrong.

Correct answer

- ☐ B Audit trails are unnecessary once a system has passed initial testing.
- ☐ C Audit trails are only useful for financial systems, not for AI-assisted decisions.
- ☐ D Audit trails slow the system down with no other benefit.

WHY

Consequential automated decisions need to be traceable and explainable after the fact — an audit trail is what makes it possible to review, justify, or correct a decision later, rather than treating the system as an unexaminable black box.

✔ **Correct** • Governance, Safety and Risk Management

A solution must keep certain categories of data within a specific country's borders due to regulatory requirements. What does this constraint require of the architecture?

- ☐ A It has no effect on architecture; data location is irrelevant to compliance.
- ☒ B **The architecture must be designed with data residency in mind, ensuring the relevant data (and its processing) stays within the required jurisdiction.** Correct answer
- ☐ C The requirement can always be safely ignored if the system is otherwise secure.
- ☐ D Data residency only applies to physical documents, never to systems like this.

WHY

Regulatory data-residency requirements are a real architectural constraint — the design has to account for where covered data is stored and processed, not treat location as an afterthought.

✔ **Correct** • Governance, Safety and Risk Management

An agent is granted broad access across many systems "just in case it's useful later," rather than only what its current task requires. What principle does this violate, and what's the risk?

- ☐ A This follows least-privilege correctly; broad access is always the safer default.
- ☒ B **It violates least-privilege; granting more access than the current task needs increases the potential damage from a bug, mistake, or misuse.** Correct answer

- ☐ C Least-privilege only applies to human users, not to automated agents.
- ☐ D There is no meaningful risk as long as the agent's prompt tells it to behave responsibly.

WHY

Least-privilege means granting only the access a task actually needs. Broad, "just in case" access widens the blast radius of any bug, mistake, or misuse well beyond what the current task requires.

✔ **Correct** • Stakeholder Communication and Lifecycle Management

A production feature depends on a specific model version that will eventually be deprecated. What lifecycle planning reduces the risk of the feature breaking unexpectedly?

- ☐ A Assuming the model version will be supported indefinitely with no planning needed.
- ☐ B Lifecycle planning is unnecessary since model versions never change.
- ☐ C Waiting until the old version is fully retired before taking any action.

☒ D **Planning ahead for model version updates and deprecations, so the transition can be tested and managed proactively rather than discovered when the old version suddenly stops working.** Correct answer

WHY

Model versions do get deprecated over time; planning for that transition ahead of time (testing against the newer version, scheduling the migration) avoids a scramble when the old version is retired unexpectedly.

✔ **Correct** • Stakeholder Communication and Lifecycle Management

After a production incident, the team holds a review meeting. What should this review focus on to be genuinely useful?

- ☐ A Closing the meeting quickly with no documented follow-up actions.
- ☒ B **Identifying process and system improvements that reduce the chance of a similar incident recurring, rather than assigning individual blame.** Correct answer

- ☐ C Avoiding any discussion of what actually went wrong.
- ☐ D Identifying which individual is to blame for the incident.

WHY

A blame-focused review tends to suppress honest information sharing and rarely fixes the underlying system issue. A review focused on process and system improvements is far more likely to prevent a repeat incident.

✔ **Correct** • Stakeholder Communication and Lifecycle Management

A proposed new use case would have a Claude-powered system process sensitive customer data in a novel way. Which of the following should happen before it's greenlit? (Select all that apply.)

- ☐ A Approval based solely on the engineering team's confidence, with no other stakeholders involved.
- ☒ B Sign-off from the business stakeholders who understand the use case's actual purpose and risk tolerance. Correct answer
- ☒ C Security review of how the sensitive data will be protected and scoped. Correct answer
- ☒ D Legal review of the data-handling and regulatory implications. Correct answer

WHY

A novel use case touching sensitive data needs legal and security review alongside informed business sign-off — engineering confidence alone, with no other stakeholders involved, skips exactly the review this kind of decision needs.

✔ **Correct** • Claude Models, Prompting and Context Engineering

How should an architect decide whether a given task needs live retrieval of external information versus relying on the model's general training knowledge?

- ☒ A Base it on whether the needed information is current, proprietary, or specific enough that the model's general training knowledge wouldn't reliably have it — retrieval matters most when that's the case. Correct answer

- ☐ B Never use retrieval under any circumstances.
- ☐ C The choice should be made randomly for each request.
- ☐ D Always use live retrieval for every task, regardless of whether the information changes or is already well-covered by the model's general knowledge.

WHY

Retrieval earns its cost and complexity when the needed information is current, proprietary, or otherwise unlikely to be reliably known from general training — for stable, well-known information, retrieval may add little value.

✔ **Correct** • Claude Models, Prompting and Context Engineering

A multi-step pipeline includes an initial triage step and a few steps that require deep reasoning. What model-tier strategy fits this well?

- ☒ A Use a fast, economical model tier for the simple triage step and reserve the most capable tier for the steps that genuinely require deep reasoning. Correct answer
- ☐ B Model tier choice should be identical across all steps regardless of what each step actually requires.
- ☐ C Use the most capable, most expensive model tier uniformly across every step, including simple triage.
- ☐ D Use the cheapest possible model tier for every step, including the deep-reasoning ones.

WHY

Matching model tier to each step's actual difficulty — economical for simple triage, capable for genuinely hard reasoning — is more cost-efficient than applying one tier uniformly across a pipeline with very different per-step demands.

✔ **Correct** • Developer Productivity and Operational Enablement

Several application teams across the org are each separately re-implementing the same Claude API integration boilerplate (auth, retries, error handling). What investment reduces this duplicated effort?

☐ A

There is no way to reduce this kind of duplicated effort.

B Building a shared internal SDK/wrapper that handles the common integration boilerplate once, for every team to reuse.

Correct answer

C Banning any team from calling the Claude API directly under any circumstances.

D Leaving every team to keep reimplementing the same boilerplate independently.

WHY

A shared internal SDK or wrapper for the common integration plumbing (auth, retries, error handling) lets every team build on the same solid foundation instead of separately reinventing it.

✔ **Correct** • Developer Productivity and Operational Enablement

Engineers frequently need production access just to diagnose why a particular Claude-powered request produced unexpected output. What investment reduces this dependency?

A Restricting all debugging information so engineers must always use production access.

B There is no way to debug these issues without direct production access.

C Removing all logging to simplify the system.

D Investing in local development and debugging tooling — such as request/response logging and replay capability — so issues can often be diagnosed without needing production access.

Correct answer

WHY

Good local tooling — logging requests/responses and enabling replay — lets engineers diagnose many issues without needing direct production access every time, improving both speed and operational safety.



Want structured training before your exam? [Explore our AI Courses](#) →

Take your AI skills further with our Online AI Courses:

[AI Courses \(FDE, GenAI & Agentic AI Engineer, ML/AI Data Science & more\)](#) →

[AI Testing Courses \(AI & LLM Testing, Claude for Automation & more\)](#) →

ISHA TRAINING SOLUTIONS

Isha Training Solutions builds hands-on AI, GenAI, Agentic AI and LLM testing courses — including this free Claude Certified Architect – Professional mock exam platform with scenario-based practice questions and full-length scored mock exams.

QUICK LINKS

[Claude Solutions Architect Course](#)

[Full CCA-F \(Architect\) Exam Prep](#)

[CCA-F Timed Mock Exams](#)

[CCAO-F \(Associate\) Exam Prep](#)

[CCDV-F \(Developer\) Exam Prep](#)

[AI Courses](#)

[Home](#)

CONTACT

 **+91-9133190573**

ishatrainingsolutions.org

Isha Training Solutions © 2026. All rights reserved. · Claude Certified Architect – Professional mock exam prep · Independent, unofficial study material — not affiliated with, endorsed by, or sponsored by Anthropic.

Built by **Isha Training Solutions**