Claude Certified Associate - Foundations (CCAO-F)

Practice Question Bank Set A

60 scenario-based items weighted to the seven CCAO-F domains. Distractors are deliberately plausible: each wrong option is a reasonable practice that does not fit the specific scenario. Includes an answer grid and an explanation for every option.

#	Domain	Exam weight	Questions here
1	Prompting and Task Execution	14%	8
2	Output Evaluation and Validation	21%	13
3	Product and Model Selection	12%	7
4	Workflow Integration and Solution Design	16%	10
5	Configuration and Knowledge Management	12%	7
6	Governance, Risk, and Responsible Use	15%	9
7	Troubleshooting and Optimization	10%	6
	Total	**100%**	**60**

How to use this set

Format mirrors the real exam: 60 items, 120 minutes, multiple choice and multiple response. When an item says **(Select TWO)** or **(Select THREE)**, you must pick exactly that many - there is no partial credit. Sit the whole set under timed conditions (about 2 minutes per item) before opening the Answer Key. Aim for 75% or better (45/60) as a comfortable margin over the 720/1000 scaled pass mark.

Unofficial practice material written to the published CCAO-F domain blueprint. These are not actual exam questions; product details change, so confirm features against Anthropic's current documentation and the official exam guide.

Questions

Domain 1: Prompting and Task Execution (14%)

Q1. A marketing coordinator asks Claude to "write a launch email for our new expense app." The draft is polished but generic, too long, and has three different calls to action. Which revision to the prompt will MOST improve the next draft?

A. Add the target audience (finance managers at mid-size firms), the single goal (book a demo), a 150-word limit, and the two product benefits to emphasise.

B. Ask Claude to critique its own draft and then rewrite it based on that critique.

C. Add a role statement such as "You are an award-winning B2B copywriter" and ask for a more compelling version.

D. Ask Claude to produce five alternative versions so the coordinator can choose the strongest one.

Q2. An analyst must turn a 40-page supplier contract into (1) a list of renewal and termination dates, (2) a risk summary, and (3) a recommendation memo for leadership. Which approach is MOST effective?

A. Paste each section of the contract into a separate chat and combine the partial outputs afterwards.

B. Ask Claude to draft the leadership memo first so the analyst can see the conclusion, then request supporting clauses.

C. Have Claude first extract the dates and key clauses with page references, verify that extraction, and then build the risk summary and memo from the verified material.

D. Ask for all three deliverables in a single detailed prompt so Claude keeps the full context in one response.

Q3. Claude is classifying 300 customer comments as Praise, Complaint, or Feature Request. Spot checks show that sarcastic comments ("Great, another update that deletes my files") are often labelled Praise. Which TWO changes will MOST directly reduce this error? **(Select TWO)**

A. Add a few labelled examples to the prompt, including sarcastic comments with their correct labels.

B. Split the 300 comments into batches of 50 so that Claude pays closer attention.

C. Ask Claude to add a confidence score to every label.

D. Add a fourth category called "Sarcasm" for ambiguous comments.

E. Define each category precisely and state that the label should reflect the customer's actual intent, not surface wording.

Q4. A team lead wants ideas for naming an internal knowledge-sharing programme. Which prompt strategy BEST suits this task?

A. Ask Claude to research naming conventions used by Fortune 500 companies, with citations, before suggesting names.

B. Give Claude a strict scoring rubric and ask it to return only names that score above 8 out of 10.

C. Ask for 15 to 20 names across several styles (playful, descriptive, acronym), each with a one-line rationale, then iterate on the shortlisted few.

D. Ask Claude for the single best name with a detailed justification, so the decision is quick.

Q5. A consultant's prompt includes background notes, a pasted client email, and instructions, all as one block of text. The client email contains the sentence "Please summarise this in one line," and Claude follows it instead of the consultant's request for a full analysis. What is the BEST fix?

A. Remove the sentence from the client email before pasting it.
B. Write the consultant's instructions in capital letters so they take precedence.
C. Move the consultant's instructions to the very top of the prompt and repeat them at the end.
D. Put the client email inside clearly labelled delimiters (for example <client_email> tags) and state that it is material to analyse, with the consultant's instructions given separately.

Q6. After two rounds of revision, a policy summary drafted by Claude is accurate but still reads as too technical for warehouse staff. What is the MOST efficient next step?

A. Ask Claude to "make it simpler" and repeat until the tone feels right.
B. Start a new chat with the original prompt so earlier revisions do not bias the output.
C. Tell Claude who the readers are, the reading level to target, which terms to replace with everyday examples, and which three rules must stay unchanged.
D. Regenerate the response several times and pick the most readable version.

Q7. A project manager wants Claude to produce a weekly status update with identical structure every week, which leadership skims in 60 seconds. Which TWO practices MOST improve consistency? **(Select TWO)**

A. Use a high-capability model for the update, since larger models follow formats perfectly.
B. Specify the required sections, their order, and a length limit for each.
C. Paste last week's update and ask Claude to "do the same for this week" without further guidance.
D. Ask Claude to choose the most suitable format each week based on that week's content.
E. Provide an example of a past update in the exact desired format.

Q8. A researcher asks Claude to compare three CRM vendors using vendor documents they have attached. Claude's comparison includes features that do not appear in any of the attachments. Which instruction BEST prevents this in future?

A. "Keep the comparison short so there is less room for error."
B. "Be accurate and do not make mistakes."
C. "Use the attached documents and your general knowledge to give the most complete comparison possible."
D. "Use only the attached documents. For each claim, cite the document and section, and write 'Not stated' when the documents do not cover a point."

Domain 2: Output Evaluation and Validation (21%)

Q9. Claude summarises a sales report and states "Region West grew 18% quarter-on-quarter." The source table shows West at $1.20M last quarter and $1.38M this quarter. What should the reviewer do?

A. Leave the figure, since the direction of the trend (growth) is correct for an executive audience.
B. Correct the figure to 15% and re-check every other calculated figure in the summary against the source table.
C. Ask Claude to re-check its maths and accept the corrected answer it gives.
D. Correct the figure to 15% and share the summary, since the remaining content is narrative.

Q10. Which feature of a Claude response about a regulation is the STRONGEST indicator of a possible hallucination?

A. It states that one requirement may differ by jurisdiction.
B. It summarises the regulation in fewer words than expected.
C. It recommends consulting a compliance professional for the final interpretation.
D. It quotes a specific clause number and wording that cannot be found in the official text of the regulation.

Q11. A clinic uses Claude to draft patient instructions that include medication-fasting guidance before procedures. What is the MOST appropriate validation step before the instructions are sent?

A. Have the practice manager proofread the instructions for clarity and spelling.
B. Review by a qualified clinician who is accountable for the medical content.
C. Ask Claude to verify each instruction against current medical guidelines and confirm accuracy.
D. Have a second AI model check the instructions and flag any disagreements.

Q12. A reviewer asks Claude the founding year of a supplier in two separate chats and gets two different answers. What is the BEST next step?

A. Use the more recent year, as it is more likely to reflect updated information.
B. Regenerate the answer several more times and use the year that appears most often.
C. Ask Claude which of the two answers is correct and why.
D. Treat the fact as unverified and confirm it from an authoritative source, such as company registry filings.

Q13. Claude drafts a list of "ideal candidate traits" for a sales role. Several traits ("young and energetic", "native English speaker") are not job-related and could disadvantage protected groups. Which TWO actions are MOST appropriate? **(Select TWO)**

A. Flag the draft as potentially biased and remove criteria that are not demonstrably job-related.
B. Replace the flagged phrases with softer synonyms such as "dynamic" and "fluent communicator".
C. Stop using Claude for recruitment content entirely.
D. Keep the list but add a statement that the company is an equal-opportunity employer.
E. Revise the prompt to ask for skills-based, measurable criteria, and have HR review the final list.

Q14. A Claude-drafted variance analysis for the CFO explains a cost overrun as "driven by a supplier price increase." The data and notes provided to Claude mention only volume changes. What should the reviewer conclude?

A. The analysis should be regenerated until Claude gives an explanation that mentions volume.
B. The explanation should be moved to an appendix so it does not affect the main analysis.
C. The explanation is unsupported by the inputs; it must be confirmed with the business owner or removed before sharing.
D. The explanation is plausible and consistent with industry trends, so it can stay with a softer phrase such as "likely driven by".

Q15. Claude produced a detailed explanation of a new IT security policy. It must now go to the IT team and to 2,000 non-technical employees. What is the BEST approach?

A. Ask Claude for one middle-ground version written for a general technical audience.
B. Send only a short employee version to everyone, since IT can read the original policy.
C. Send the detailed version to everyone, since a single source of truth avoids confusion.
D. Create two adapted versions: an implementation-focused version for IT and a short action-focused version for employees. Then check both against the policy text.

Q16. A user wants a two-page onboarding checklist that they will refine over several turns, and then share and download. Which output format is MOST appropriate?

A. An artifact.
B. Inline text in the chat response.
C. A bulleted summary inside a Project instruction.
D. A JSON object with one key per checklist item.

Q17. An accounts-payable team wants Claude to extract vendor, invoice date, amount, and PO number from scanned invoices so the data can be imported into their ERP system. Which output specification is BEST?

A. An artifact containing a formatted report of all invoices.
B. A fixed set of named fields in a consistent structured format (for example CSV), with a defined value for missing data and a fixed date format.
C. A clear table per invoice, with column names chosen by Claude to suit each invoice.
D. A short paragraph per invoice highlighting the key details.

Q18. Which TWO of the following outputs MOST clearly require human review before use? **(Select TWO)**

A. A draft reply to a formal information request from a financial regulator.
B. A customer refund approval that exceeds the support agent's delegated authority.
C. A list of icebreaker questions for an internal team meeting.
D. Five alternative subject lines for an internal newsletter.
E. A summary of a public industry article for the user's own background reading.

Q19. A reviewer is checking a Claude-generated summary of a 30-page research report. Which technique is MOST effective at detecting important OMISSIONS?

A. Check that the summary's length is roughly 10% of the original.
B. Map the summary against the report's section structure and stated key findings to confirm each major point is represented.
C. Spot-check three figures in the summary against the report.
D. Ask Claude whether anything important was left out of the summary.

Q20. Claude provides a market-size estimate and adds: "This is based on information up to my training cutoff and may be outdated." The estimate will go into a board paper. What is the BEST handling?

A. Verify the figure against a current authoritative source, or use a web-enabled research feature, and cite that source in the board paper.
B. Keep the figure and the caveat word-for-word in the board paper so readers can judge for themselves.
C. Discard the entire analysis because one figure may be outdated.
D. Remove the caveat but round the figure to signal it is approximate.

Q21. Claude's vendor comparison table lists Vendor A support as "24/7". Vendor A's proposal says "business-hours support; 24/7 available as a paid add-on." How should this be classified and handled?

A. As oversimplification that misrepresents the source: correct the cell to show the condition and cost, then re-check other cells for similar lost nuance.
B. As a hallucination: regenerate the whole comparison from scratch.
C. As a formatting limitation of tables: add a footnote to the table and leave the cell as it is.
D. As acceptable summarisation, since 24/7 support is technically available from Vendor A.

Domain 3: Product and Model Selection (12%)

Q22. A support operation must route about 5,000 short tickets per hour into 8 categories. Accuracy on this simple task is already high with any model, and cost and latency are the main concerns. Which model family is the BEST fit?

A. Alternate between Sonnet and Opus to balance the load
B. Sonnet
C. Haiku
D. Opus

Q23. A strategy director wants Claude to reason through a complex acquisition scenario across many long documents, with subtle trade-offs. Quality matters most; cost and speed are secondary. Which model family is MOST appropriate?

A. Sonnet, because it is the default for business work
B. Haiku
C. Opus
D. Haiku, with the documents summarised first to save cost

Q24. An operations team needs a model for everyday drafting, summarisation, and analysis across many users. They want strong quality but must keep costs reasonable at scale. Which model family is the MOST sensible default?

A. Rotate models weekly and compare
B. Haiku for every request, to minimise cost
C. Opus for every request, to guarantee quality
D. Sonnet

Q25. A consultant will work on the same client engagement for three months. Each week they reuse the same 12 background documents and want Claude to follow the same writing conventions. Which approach BEST fits?

A. Save the documents in an artifact and refer to it in each new chat.
B. Start a new chat each week and upload the 12 documents again.
C. Use one very long chat for the whole engagement so Claude remembers everything.
D. Create a Project with custom instructions for the conventions and the 12 documents in project knowledge.

Q26. An analyst needs a briefing on competitor announcements from the past two weeks, drawn from many sources, with citations. Which approach is MOST appropriate?

A. Use research mode, which searches across multiple sources and returns a cited report.
B. Add last quarter's competitor reports to a Project and ask Claude to summarise recent announcements.
C. Ask in a standard chat without web access and request citations.
D. Ask Claude to produce the briefing as an artifact so it can be updated later.

Q27. Over a very long conversation, Claude has begun ignoring budget constraints agreed near the start, and responses are slower. What is the BEST action?

A. Restate the budget constraints in the next message and continue in the same chat.
B. Switch to a smaller, faster model to reduce latency in the current chat.
C. Edit the first message to emphasise the constraints and regenerate the latest response.
D. Ask Claude to summarise the current state, decisions, and constraints, check the summary, and continue in a fresh chat (or save it to the Project's knowledge).

Q28. Which TWO factors should MOST directly drive the choice between Haiku, Sonnet, and Opus for a new workflow? **(Select TWO)**

A. Volume, latency, and cost constraints.

B. The length of the user's prompt.

C. Which model is the newest release, since newer is always better.

D. Whether the output will be shown as an artifact or inline.

E. The complexity of the task and the quality required.

Domain 4: Workflow Integration and Solution Design (16%)

Q29. HR wants to use Claude to speed up job-description drafting. Which workflow design is BEST?

A. Hiring managers prompt Claude directly and publish the result, with HR auditing a sample each quarter.

B. Each recruiter builds their own prompt, and the best prompt is chosen after six months.

C. Hiring managers complete a standard intake form, Claude drafts from it in a configured Project, and HR reviews for accuracy, inclusive language, and compliance before publishing.

D. Claude drafts from past job descriptions alone, and HR publishes if the format matches the template.

Q30. A business analyst has 15 sets of stakeholder interview notes for a new expense-approval process. How can Claude add the MOST value at this stage?

A. Generate additional requirements that stakeholders may have forgotten and add them to the list.

B. Synthesise the notes into themes, conflicting requirements, gaps, and open questions for stakeholders to resolve.

C. Draft the final requirements specification and send it for sign-off.

D. Recommend which stakeholder's requirements should take priority when they conflict.

Q31. A director asks for a one-paragraph explanation of what Claude can realistically do for their customer-service team. Which statement is MOST accurate?

A. "Claude can draft replies, summarise case history, and suggest next steps to speed up agents. Drafts need agent review, Claude sees account data only through connections we configure, and escalation decisions stay with staff."

B. "Claude automatically has secure access to all our systems, so agents will no longer need to look things up."

C. "Claude can handle most tickets end-to-end, and agents only need to step in when a customer complains."

D. "Claude can draft replies but cannot summarise or analyse information, so its value is limited to typing speed."

Q32. A process lead wants to use Claude to improve a 12-step invoice-processing workflow. What is the BEST first use of Claude?

A. Analyse the current process documentation and staff interviews to map the steps, identify bottlenecks and repetitive manual work, and propose where AI assistance fits while keeping controls.

B. Draft a new 6-step process that removes approvals Claude judges redundant.

C. Automate the three slowest steps immediately, then measure the impact.

D. Benchmark the workflow against industry best practice and adopt that standard.

Q33. Which TWO characteristics MOST strongly indicate that a task is a good candidate for Claude assistance? **(Select TWO)**

A. Drafts can be reviewed and corrected before they have any external effect.

B. The task is high-volume, but no one can tell what a good output looks like.

C. The task requires an immediate, legally binding decision with no human involvement.

D. The task depends on real-time data that Claude cannot be connected to.

E. The task is repetitive and language-heavy, with clear quality criteria a reviewer can check.

Q34. A training manager used Claude to design an onboarding curriculum. Pilot trainees report that module 3 (expense systems) is confusing. What reflects the BEST iterative approach?

A. Add a longer appendix to module 3 with more detail.

B. Wait for more cohorts' feedback before changing anything, so changes are based on more data.

C. Regenerate the full curriculum with a better prompt, since the original design may have other hidden issues.

D. Give Claude the specific feedback and module 3 content, request targeted revisions with worked examples, then re-test with the next cohort.

Q35. A sales team wants Claude to prepare one-page pre-meeting client briefs. Which design BEST balances efficiency and accuracy?

A. Claude writes the brief from its general knowledge of the client's industry, to save connector setup time.

B. Claude writes the brief and emails it directly to the client as a meeting agenda.

C. Claude summarises CRM notes and recent emails available through approved connectors, and the salesperson verifies key facts before the meeting.

D. Claude writes the brief, and a sales manager approves every brief before the salesperson sees it.

Q36. A department's monthly report takes 3 days, mostly spent consolidating data and writing commentary. The head asks whether to augment or fully redesign the process with Claude. What is the BEST recommendation?

A. Use Claude to estimate missing figures so consolidation can be skipped.

B. Fully redesign now, since the consolidation and commentary steps are both clearly suited to AI.

C. Keep the process unchanged until the data warehouse is modernised.

D. Start by augmenting, with Claude drafting commentary from the consolidated data under analyst review. Measure time saved and error rates, then decide on deeper redesign.

Q37. A product owner uses Claude to turn rough meeting notes into user stories, and wants the whole team to do the same consistently. What is the BEST approach?

A. Have one team member generate all user stories, so only one person uses Claude.

B. Create a shared Project whose instructions define the story format and acceptance-criteria standard, with an example story in knowledge, and keep product-owner review before stories enter the backlog.

C. Share a recommended prompt on the team wiki and let each person adapt it.

D. Let Claude add stories directly to the backlog to save time, with fixes made during sprint planning.

Q38. Which metric would MOST convincingly show stakeholders the value of a Claude-assisted contract-review workflow?

A. Reduction in average review time per contract, with QA sampling showing no increase in missed issues.
B. Average length of Claude's contract summaries.
C. Percentage of staff who say they like using Claude.
D. Number of contracts processed through Claude per month.

Domain 5: Configuration and Knowledge Management (12%)

Q39. A legal-operations team configures a Claude Project for first-pass NDA review. Which set of Project instructions is MOST effective?

A. "Review NDAs. Be concise." Users then add their own preferences in each chat.
B. Role and purpose; compare each NDA against the attached standard positions; output a table of clause, deviation, risk level, and suggested fallback; flag unclear items for attorney review; never state that an NDA is approved.
C. A full copy of the company's 40-page legal policy manual pasted into the instructions.
D. "You are an expert contracts lawyer. Review NDAs thoroughly and give your professional opinion."

Q40. A Project's knowledge contains both the 2024 and the 2026 travel policies. Claude sometimes quotes the outdated 2024 meal limits. What is the BEST fix?

A. Tell users to double-check any meal limit Claude quotes.
B. Add an instruction to "always use the latest information" and leave the knowledge as it is.
C. Upload a second copy of the 2026 policy so it carries more weight.
D. Remove the 2024 policy from project knowledge, or if it must be kept, label it clearly as superseded and instruct Claude to use only the current policy.

Q41. A team wants Claude to reference operating procedures stored in Google Drive that are updated weekly by different owners. Which approach BEST keeps answers current?

A. Upload copies to project knowledge and assign someone to re-upload them monthly.
B. Paste the relevant procedure into the chat each time it is needed.
C. Use the Google Drive connector so Claude can access current versions within each user's existing permissions.
D. Export all procedures into one consolidated PDF, upload it, and refresh it quarterly.

Q42. Which TWO practices BEST support ongoing maintenance of a team's shared Claude Project? **(Select TWO)**

A. Assign a named owner who periodically reviews instructions and knowledge files for accuracy.
B. Freeze the instructions after launch so behaviour stays predictable.
C. Allow any member to edit the instructions at any time without telling anyone.
D. Update instructions and knowledge when processes change or recurring output problems appear, and keep a short change log.
E. Upload every potentially related document so Claude never lacks information.

Q43. A Project's instructions say "Always respond in under 100 words," but users regularly ask for detailed two-page reports in the same Project. What is the BEST configuration change?

A. Remove the length rule entirely so Claude decides the length each time.
B. Revise the instructions to set the short default but allow longer outputs when a user explicitly asks for a detailed deliverable, or set up a separate Project for long-form reports.
C. Ask users to begin each long request with "Ignore your instructions".
D. Keep the rule and train users to stop asking for long reports in this Project.

Q44. When adding reference material to a team Project, which practice MOST improves answer quality?

A. Add as many documents as the limit allows so Claude has maximum context.
B. Add only a summary of each document written by Claude, to save space.
C. Add documents as photos of printed pages, to preserve their original layout.
D. Add focused, current, clearly titled documents relevant to the Project's purpose, and remove drafts and duplicates.

Q45. After a user connects Gmail to Claude, a colleague worries that Claude can now read every employee's mailbox. Which response is MOST accurate?

A. Connectors can only read email subject lines, never message content.
B. Connectors grant organisation-wide access, which is why only admins should use them.
C. The connector works within the connecting user's own authorised access and the organisation's admin controls; it does not grant access beyond that.
D. The connector can read any mailbox the user has ever emailed.

Domain 6: Governance, Risk, and Responsible Use (15%)

Q46. An employee wants to paste a spreadsheet of customer names, government ID numbers, and account balances into their personal Claude account to build a quick summary. What is the MOST appropriate action?

A. Proceed, but delete the conversation immediately afterwards.
B. Do not proceed. Use only organisation-approved tools and follow the data-classification policy, removing or anonymising identifiers where the policy requires it.
C. Proceed, but remove the names and keep the ID numbers, which are just numbers.
D. Proceed, since the summary will be used internally only.

Q47. Which use of Claude is LEAST appropriate without significant human oversight?

A. Automatically issuing final denials of insurance claims.
B. Drafting an internal FAQ from approved policy documents.
C. Summarising meeting notes for attendees.
D. Suggesting agenda items for a team offsite.

Q48. Company policy requires disclosure when AI materially contributes to client deliverables. A consultant used Claude to draft most of a client report and then edited it. What should they do?

A. Skip disclosure, since substantial human editing makes the report the consultant's own work.
B. Disclose only if the client asks how the report was produced.
C. Rewrite the report so it no longer reads like AI output, and then skip disclosure.
D. Disclose AI assistance as the policy requires, and confirm they personally reviewed and validated the content.

Q49. HR wants to use Claude to analyse 800 open-text comments from an anonymous employee survey. Which TWO practices BEST protect respondents? **(Select TWO)**

A. Share the raw comments with all managers along with Claude's analysis, for transparency.
B. Report themes at an aggregate level and avoid quoting comments that could identify individuals.
C. Attach each respondent's department and tenure so the analysis is more precise.
D. Remove names, employee IDs, and other direct identifiers before analysis.
E. Ask Claude to infer which team each comment came from, to target interventions.

Q50. A manager asks Claude about legal requirements for terminating an employee in a jurisdiction with strict labour laws. How should the manager treat Claude's response?

A. As useful background for preparing questions; the decision itself must be reviewed by qualified HR or legal counsel for that jurisdiction.
B. As irrelevant, since Claude is not able to discuss employment law.
C. As reliable if Claude is asked to confirm its answer twice.
D. As reliable guidance to act on, provided it includes citations to the relevant law.

Q51. A colleague suggests installing an unapproved browser extension that sends Claude chat content to a third-party service "to save time on formatting." What is the BEST response?

A. Don't use it. Submit the extension through the organisation's tool-approval process so its security and data handling can be assessed.
B. Use it only for chats that contain no customer data.
C. Use it for a trial week and report the results to IT afterwards.
D. Use it, since Claude itself is approved and the extension only handles formatting.

Q52. When Claude assists with a decision that affects customers, where does accountability for the decision rest?

A. Shared equally between the organisation and the model.
B. With the people and organisation that made and acted on the decision.
C. With no one, if the decision followed the AI's recommendation exactly.
D. With the AI vendor, since the recommendation came from its model.

Q53. Marketing asks Claude to write five customer testimonials for the website "in the style of happy customers," because the team has not collected real ones yet. What is the correct assessment?

A. Acceptable if each testimonial is realistic and reflects the product's actual features.
B. Acceptable if the testimonials use first names only.
C. Acceptable as placeholders that go live until real testimonials are collected.
D. Inappropriate: presenting fabricated testimonials as genuine is deceptive and may breach consumer-protection law and company policy.

Q54. A financial-services adviser wants to use Claude to draft client communications about investment products. Which consideration is MOST specific to this regulated context?

A. Communications should be sent promptly after drafting.
B. Communications should be personalised with the client's first name.
C. Communications must meet sector rules (for example fair and balanced presentation, required disclosures, and record-keeping) and pass the firm's compliance review.
D. Communications should be drafted in the firm's brand voice.

Domain 7: Troubleshooting and Optimization (10%)

Q55. Claude's summaries keep giving general industry commentary instead of summarising the attached report. What should be checked FIRST?

A. Whether the report is too long and needs to be split into chapters.
B. Whether the response length limit is too low.
C. Whether a more capable model would follow the request better.
D. Whether the prompt explicitly says to summarise the attached document, and whether the file was actually read (for example, by asking Claude to quote its opening section).

Q56. Project instructions say "Be extremely concise," while a user's message asks for "a thorough, detailed analysis." Output length is inconsistent from one run to the next. What is the BEST diagnosis and fix?

A. A model capability limit. Switch to a larger model.
B. Random model variability. Regenerate until the desired length appears.
C. Conflicting instructions. Revise the Project instructions to set precedence, for example "default to concise unless the user explicitly asks for detail."
D. User error. Ask users not to request detail in this Project.

Q57. Analysts spend about 15 minutes in every new chat re-explaining the same report-formatting rules. Which optimisation is MOST effective?

A. Drop the formatting rules and fix formatting manually afterwards.
B. Put the formatting rules and a sample report in a shared Project's instructions and knowledge.
C. Keep a text file of the rules and paste it at the start of each chat.
D. Use a single long-running chat so the rules are only explained once.

Q58. A data-extraction prompt is correct on about 85% of invoices. Most errors occur on invoices that use DD/MM/YYYY dates instead of MM/DD/YYYY. Which TWO actions BEST address this? **(Select TWO)**

A. Ask Claude to be more careful with dates.
B. Specify the required output format (YYYY-MM-DD) and a rule for ambiguous dates, such as using the vendor's country or flagging for review.
C. Add examples of invoices in both date formats with the expected normalised output.
D. Remove the date field and have staff enter dates manually.
E. Switch to the largest available model, since extraction errors indicate a capability limit.

Q59. A team changed a prompt and believes output quality improved. What is the BEST way to confirm this?

A. Ask the team to vote on which version they prefer.
B. Run the new prompt on a few recent inputs and check the outputs look good.
C. Ask Claude to compare the two prompts and say which is better.
D. Run the old and new prompts on the same representative set of inputs and score both against defined quality criteria.

Q60. A user asks Claude for a complete 30-page procedures manual in one request. The output stops partway through, and the later sections are thin. What is the BEST approach?

A. Ask Claude to compress the manual into 10 pages so that it fits.
B. Switch to a faster model to finish the output before it stops.
C. Agree an outline first, then generate and review the manual section by section, using the outline as the reference.
D. Regenerate the same request until a complete version appears.

Answer Grid

Q1	**A**	Q2	**C**	Q3	**AE**	Q4	**C**	Q5	**D**	Q6	**C**
Q7	**BE**	Q8	**D**	Q9	**B**	Q10	**D**	Q11	**B**	Q12	**D**
Q13	**AE**	Q14	**C**	Q15	**D**	Q16	**A**	Q17	**B**	Q18	**AB**
Q19	**B**	Q20	**A**	Q21	**A**	Q22	**C**	Q23	**C**	Q24	**D**
Q25	**D**	Q26	**A**	Q27	**D**	Q28	**AE**	Q29	**C**	Q30	**B**
Q31	**A**	Q32	**A**	Q33	**AE**	Q34	**D**	Q35	**C**	Q36	**D**
Q37	**B**	Q38	**A**	Q39	**B**	Q40	**D**	Q41	**C**	Q42	**AD**
Q43	**B**	Q44	**D**	Q45	**C**	Q46	**B**	Q47	**A**	Q48	**D**
Q49	**BD**	Q50	**A**	Q51	**A**	Q52	**B**	Q53	**D**	Q54	**C**
Q55	**D**	Q56	**C**	Q57	**B**	Q58	**BC**	Q59	**D**	Q60	**C**

Answer Key with Explanations

Domain 1: Prompting and Task Execution

Q1. Answer: A

The draft failed because Claude lacked context and constraints: audience, goal, length, and key content. A supplies exactly those. C can shift tone, but a role alone does not supply the missing facts or constraints. D produces more drafts from the same thin brief. B can polish wording, but Claude still does not know the audience or the single CTA the coordinator needs.

Q2. Answer: C

Decomposing the task into sequential, checkable steps (C) stops an extraction error from silently spreading into the risk summary and memo. D is not wrong in principle, but it mixes extraction, analysis, and persuasion, which makes errors hard to isolate. B anchors on a conclusion before the facts are confirmed. A breaks cross-references, since contract clauses often depend on definitions found elsewhere in the document.

Q3. Answer: A, E

Few-shot examples of the failure case (A) and precise definitions based on intent (E) target the root cause. C may help with triage, but it does not change how sarcasm is interpreted. B can help with very large inputs, but batch size is not causing this error. D breaks the required three-category taxonomy, because a sarcastic comment is still a Complaint.

Q4. Answer: C

Brainstorming benefits from breadth across distinct directions, followed by iteration (C). D narrows to one option too early. A applies a research-task approach to a creative task and slows it down without improving the ideas. B filters the ideas before the team has seen the range, using a subjective self-score.

Q5. Answer: D

The problem is that Claude cannot tell content apart from instructions. Delimiting the content and labelling its role (D) solves this structurally for any pasted material. C can help a little, but the boundary is still ambiguous. B adds emphasis, not structure. A works for this one email but does not scale and alters the source material.

Q6. Answer: C

Specific, targeted feedback about the audience, the level, what to change, and what to keep (C) closes the gap in the fewest iterations. D relies on chance. A gives direction but not enough of it, so it tends to take many cycles and risks losing required content. B throws away revisions that are already accurate.

Q7. Answer: B, E

An exemplar (E) and explicit structural constraints (B) are the main levers for repeatable format. D adds variability on purpose. A may improve quality, but model choice does not replace a specified format. C is close, but without stated rules Claude may copy last week's content patterns or drift in structure.

Q8. Answer: D

Restricting the comparison to the sources, requiring citations, and explicitly permitting "Not stated" (D) removes the pressure to fill gaps and makes every claim checkable. B states a goal but gives Claude no mechanism for meeting it. C invites exactly the unsupported claims the researcher saw. A reduces volume, not unsupported content.

Domain 2: Output Evaluation and Validation

Q9. Answer: B

(1.38 - 1.20) / 1.20 = 15%. One calculation error is evidence that other derived figures may be wrong too, so the reviewer should verify all of them (B). D fixes the known error but assumes there are no others. C relies on self-correction instead of independent verification. A accepts a material misstatement.

Q10. Answer: D
Specific, confident details that cannot be traced to the authoritative source (D) are the classic pattern of a hallucination. A and C are appropriate caveats. B is about brevity, not fabrication.

Q11. Answer: B
Patient-safety content needs review by a qualified, accountable expert (B). C and D may catch some errors, but neither provides professional accountability. A checks readability, not clinical accuracy.

Q12. Answer: D
Inconsistency between runs signals low reliability for that fact, so it needs external verification (D). B can converge confidently on a wrong value. C asks the same unreliable source to judge itself. A is an assumption with no basis.

Q13. Answer: A, E
Recognising the bias and removing criteria that are not job-related (A), then correcting the prompt and adding human review (E), is the right response. D leaves discriminatory criteria in place. B can hide the same bias behind coded language, although "fluent communicator" may be legitimate if the job truly requires it. C overreacts, because properly reviewed use is acceptable.

Q14. Answer: C
A causal claim that the provided material does not support is a fabrication risk, whatever its plausibility (C). D hedges the claim but still presents an unverified cause to the CFO. B still circulates it. A tries to steer the output toward an expected answer instead of validating it.

Q15. Answer: D
Adapting content to each audience while checking it stays faithful to the source (D) serves both groups. C overloads non-technical staff. B under-serves IT. A compromises in a way that serves neither audience well.

Q16. Answer: A
Substantial, standalone content that will be iterated on and reused fits an artifact (A). Inline text (B) suits short conversational answers and is harder to maintain across revisions. JSON (D) is designed for machines, not people. Project instructions (C) configure Claude's behaviour; they are not a place for deliverables.

Q17. Answer: B
Automated import needs strict, predictable structure, including rules for missing values and formats (B). C looks structured, but varying column names will break the import. D and A are readable by people but require manual re-keying.

Q18. Answer: A, B
Regulatory correspondence (A) and decisions beyond delegated authority (B) carry legal, financial, and accountability risk. C, E, and D are low-stakes, internal, and easily corrected.

Q19. Answer: B
To find what is missing, compare coverage against the source's structure (B). C detects errors in what is present, not omissions. D is a useful supplementary prompt, but it is not independent verification. A measures length, not coverage.

Q20. Answer: A
The caveat points to exactly what needs verifying. A board paper needs a current, citable figure (A). B passes unresolved uncertainty to the board. D hides the uncertainty. C discards analysis that may otherwise be useful.

Q21. Answer: A
Turning a conditional, paid option into an unconditional feature misrepresents the source (A), and the same pattern often affects other cells. D accepts the distortion. C leaves the misleading cell in place. B mislabels the error, since the information exists, and regenerating does not guarantee a fix.

Domain 3: Product and Model Selection

Q22. Answer: C
For a simple, high-volume task where cost and latency dominate, the fastest, most cost-efficient family (C) fits best. B would work but costs more without a needed gain. D is the most expensive and is overkill. A adds cost and inconsistency.

Q23. Answer: C
The most capable family (C) suits complex, high-stakes reasoning when quality is the priority. Sonnet (A) is a strong general default, but the scenario explicitly puts quality above everything else. B and D optimise for speed and cost, the opposite of the stated priorities.

Q24. Answer: D
The balanced family (D) is the usual default for broad knowledge work, combining strong capability with manageable cost and speed. C over-spends on routine tasks. B may under-deliver on analysis. A is not a deployment strategy.

Q25. Answer: D
Projects provide persistent instructions and knowledge across many chats (D). C will eventually run into context limits and degrade. B is repetitive and prone to inconsistency. Artifacts (A) are outputs, not a reusable knowledge base for new chats.

Q26. Answer: A
Current, multi-source, cited information needs research or web capabilities (A). C cannot reach recent events and may produce unverifiable citations. B uses outdated material. D only changes the output format; it does nothing to gather information.

Q27. Answer: D
Long chats dilute or exceed the effective context. A checked summary carried into a fresh context fixes the root cause (D). A is a reasonable short-term patch, but the chat keeps growing and the problem will come back. B trades quality for speed without fixing the problem. C may discard all later conversation from that point.

Q28. Answer: A, E
Model choice balances task complexity and quality needs (E) against volume, speed, and budget (A). C confuses release date with fit for the task. D is a formatting decision. B is not a primary factor in itself, although a very large context may be.

Domain 4: Workflow Integration and Solution Design

Q29. Answer: C
A good design defines standard inputs, a configured AI step, an accountable review, and a clear handoff (C). A reviews only after publication, so errors reach candidates first. D skips role-specific inputs and checks format rather than content. B is inconsistent and delays standardisation.

Q30. Answer: B
Synthesis that surfaces conflicts and gaps for humans to resolve (B) uses Claude's strengths and keeps decision rights with stakeholders. C skips the resolution step. D hands Claude a business-priority decision that belongs to people. A adds requirements nobody stated, without validating them.

Q31. Answer: A
A balances value with real limitations (review, data access, accountability). C overstates autonomy and reliability. D understates Claude's capability. B is false, because access depends on configured connectors and permissions.

Q32. Answer: A
Understanding the current state before changing it (A) targets real problems and keeps controls in place. C acts before diagnosis. B removes controls based on AI judgement without finance governance. D may be useful later, but generic benchmarks ignore this organisation's specific constraints.

Q33. Answer: A, E
Language-heavy, repeatable work with checkable criteria (E) and a review point before impact (A) are strong indicators. C removes accountability. D is a data-access mismatch. B makes validation impossible, whatever the volume.

Q34. Answer: D
Targeted revision based on specific evidence, then re-testing (D), is efficient iteration. C discards modules that worked. A adds length without addressing the confusion. B delays a fix for a clearly identified problem.

Q35. Answer: C
Grounding Claude in approved, relevant data and adding a lightweight verification by the owner (C) is efficient and accurate. A is not grounded in client-specific data. B sends unreviewed content to an external party. D adds a heavy approval bottleneck to a low-risk internal document.

Q36. Answer: D
An evidence-based, incremental rollout (D) proves value while controlling risk. B commits to large change without evidence. C delays value that is available now. A replaces real data with estimates and undermines data integrity.

Q37. Answer: B
Shared configuration plus review gives consistency and quality at team scale (B). C is a reasonable start, but individual adaptations cause drift. A creates a bottleneck. D removes the quality gate.

Q38. Answer: A
Value means efficiency gains while quality is maintained (A). D measures activity, not value. C is a useful adoption signal but not proof of value. B says nothing about outcomes.

Domain 5: Configuration and Knowledge Management

Q39. Answer: B
Effective instructions define the role, the reference material, the output structure, and the boundaries (B). D is vague and invites Claude to act as an authority. C belongs in knowledge, not instructions, and buries the actual directives. A leaves format and standards to each user.

Q40. Answer: D
Conflicting, stale knowledge should be curated at the source (D). C adds clutter and is unreliable. B is vague, because Claude may not reliably tell which document is newer. A pushes the fix onto every user instead of correcting the configuration.

Q41. Answer: C
A connector provides access to live sources and respects existing permissions (C). A and D leave gaps where the knowledge is stale between refreshes. B is manual and error-prone, although it works for one-off use.

Q42. Answer: A, D
Ownership with periodic review (A) and controlled, documented updates (D) keep a Project reliable. B lets the Project go stale. E adds noise and conflicting material. C causes untracked drift.

Q43. Answer: B
Instructions should reflect real use, with defaults plus explicit exceptions or Projects separated by purpose (B). D ignores a legitimate need. A loses a useful default. C is a fragile workaround that weakens the configuration.

Q44. Answer: D
Curated, relevant, clearly labelled knowledge improves retrieval and reduces conflicts (D). A dilutes the context with irrelevant material. B loses detail and may introduce summary errors. C can reduce extraction quality.

Q45. Answer: C
Connectors act under the authenticated user's permissions and the organisation's settings (C). B and D overstate the scope. A understates what a connector can do.

Domain 6: Governance, Risk, and Responsible Use

Q46. Answer: B
Sensitive personal data must stay within approved tools and data policy (B). Deleting the chat afterwards (A) does not undo the disclosure. ID numbers (C) are themselves sensitive identifiers. Internal use (D) does not make an unapproved tool acceptable.

Q47. Answer: A
Consequential decisions that affect people's rights and finances (A) need accountable human decision-makers. The other options are low-risk drafting tasks that are easy to review.

Q48. Answer: D
Follow the organisation's disclosure policy and keep accountable review (D). A, B, and C each reinterpret or circumvent the policy.

Q49. Answer: B, D
De-identification (D) and aggregate reporting (B) protect anonymity. C and E add or infer identifying details. A exposes raw, potentially identifiable comments.

Q50. Answer: A
Claude can help someone understand the issues, but a high-consequence regulated decision needs qualified professionals (A). Citations (D) and repeated confirmation (C) do not create professional accountability. B understates Claude's legitimate usefulness.

Q51. Answer: A
Unapproved tools that transmit data to third parties need governance review first (A). B relies on each person's judgement about what is sensitive. C and D bypass the process.

Q52. Answer: B
AI is a tool, and accountability remains with the humans and organisation (B).

Q53. Answer: D
Fabricated endorsements presented as real are deceptive (D). Names (B), realism (A), or a temporary status (C) do not change that, and "placeholders" that go live are published content.

Q54. Answer: C
Regulated sectors impose specific content, disclosure, and record-keeping obligations (C). B, D, and A are general good practice, not regulatory considerations.

Domain 7: Troubleshooting and Optimization

Q55. Answer: D
Start with the most likely and cheapest causes: an ambiguous instruction or input that was not read (D). C may help eventually, but changing models before checking the inputs skips the diagnosis. A is possible, but only after confirming the file was read at all. B is unrelated to the content going off topic.

Q56. Answer: C
Inconsistency between two competing instructions is a configuration conflict. Explicit precedence resolves it (C). B and A misdiagnose the problem. D blocks a legitimate need.

Q57. Answer: B
Persisting recurring context in a Project removes repeated setup and standardises output (B). C helps but stays manual and prone to error. D eventually runs into context-length problems. A moves the effort instead of removing it.

Q58. Answer: B, C
Examples that target the failure pattern (C) and explicit format and ambiguity rules (B) directly fix the diagnosed cause. A is too vague to change behaviour reliably. E may help marginally but ignores the specific, fixable cause. D gives up the automation.

Q59. Answer: D
A controlled comparison on representative inputs with defined criteria produces evidence (D). C is not grounded in output data. B has no baseline to compare against. A measures opinion, not quality against criteria.

Q60. Answer: C
Long deliverables are more reliable and easier to review when produced from an outline, section by section (C). D repeats the same limit. A changes the requirement. B misdiagnoses the problem as speed.