Claude Certified Associate - Foundations (CCAO-F)

Practice Question Bank Set B

A second, fully independent set of 60 scenario-based items with no overlap with Set A, weighted to the seven CCAO-F domains. Distractors are deliberately plausible. Includes an answer grid and an explanation for every option.

#	Domain	Exam weight	Questions here
1	Prompting and Task Execution	14%	8
2	Output Evaluation and Validation	21%	13
3	Product and Model Selection	12%	7
4	Workflow Integration and Solution Design	16%	10
5	Configuration and Knowledge Management	12%	7
6	Governance, Risk, and Responsible Use	15%	9
7	Troubleshooting and Optimization	10%	6
	Total	**100%**	**60**

How to use this set

Format mirrors the real exam: 60 items, 120 minutes, multiple choice and multiple response. When an item says **(Select TWO)** or **(Select THREE)**, you must pick exactly that many - there is no partial credit. Sit the whole set under timed conditions (about 2 minutes per item) before opening the Answer Key. Aim for 75% or better (45/60) as a comfortable margin over the 720/1000 scaled pass mark.

Unofficial practice material written to the published CCAO-F domain blueprint. These are not actual exam questions; product details change, so confirm features against Anthropic's current documentation and the official exam guide.

Questions

Domain 1: Prompting and Task Execution (14%)

Q1. A procurement officer asks Claude: "Analyse these three supplier quotes." The response describes each quote at length but does not help with the decision. What is the MOST important information missing from the prompt?

A. An instruction to use a table instead of paragraphs.
B. The decision to be made and the criteria that matter (for example total 3-year cost, delivery lead time, and warranty), plus the desired output (a ranked recommendation with rationale).
C. A request for Claude to be thorough and consider every aspect of the quotes.
D. The names and job titles of the people who will read the analysis.

Q2. A communications manager needs a crisis holding statement. They have the facts, the approved key messages, and legal's list of phrases to avoid. Which prompt structure is BEST?

A. Describe the situation briefly and ask Claude to draft a statement "in line with crisis communications best practice".
B. Ask Claude to write the statement first, then paste in legal's list and ask it to fix any problems.
C. Give the facts, key messages, and prohibited phrases as clearly labelled sections, then state the task, the audience, the tone, and a word limit.
D. Paste in a previous crisis statement and ask Claude to update the details.

Q3. A team asks Claude to "improve" a 15-slide sales deck's speaker notes. The result changes facts and product claims along with the wording. Which TWO prompt changes would BEST prevent this? **(Select TWO)**

A. Define "improve" with concrete criteria, such as a maximum of 90 seconds per slide and a plain-language register.
B. Ask Claude to explain each change after making it.
C. Ask Claude to be more conservative in its edits.
D. Ask Claude to rewrite the notes from scratch based on the slide titles alone.
E. State explicitly what to change (clarity, flow, and length) and what must not change (figures, product claims, and pricing).

Q4. An HR generalist wants Claude to draft a difficult performance-improvement conversation script. Which approach is MOST effective?

A. Ask for a script that avoids mentioning specific performance issues, to keep the conversation positive.
B. Provide the specific performance gaps with examples, the desired outcome, the company's PIP policy points, and the tone (direct but supportive), then ask for a script with likely employee responses and suggested replies.
C. Ask Claude for a generic PIP conversation template and adapt it by hand.
D. Ask Claude to role-play the employee first, so the HR generalist can practise before any script exists.

Q5. Claude is asked to "research, outline, and write" a 3,000-word thought-leadership article in one prompt. The result is shallow and the structure is weak. What is the BEST way to restructure the task?

A. Increase the requested length to 5,000 words so Claude covers each topic in more depth.

B. Add "think step by step" to the original single prompt.

C. Ask for three complete article drafts and merge the best paragraphs.

D. Split it into stages: research summary, then an agreed outline with key arguments, then drafting section by section, with review between stages.

Q6. A user wants Claude's meeting summaries to look like the team's existing ones: a decision log, owner-tagged action items, and open questions. Which technique MOST reliably achieves this?

A. Ask Claude to use the most common professional meeting-summary format.

B. Ask Claude to output JSON first and convert it to prose later.

C. Describe the format in a single sentence and trust Claude to infer the rest.

D. Include one or two real example summaries in the desired format, along with a brief description of each section.

Q7. Which TWO practices BEST illustrate adapting the prompting strategy to the task type? **(Select TWO)**

A. For brainstorming, require Claude to cite sources for each idea.

B. For drafting, specify audience, purpose, tone, and length, and provide any must-include points.

C. For every task, use the same detailed prompt template regardless of the type of work.

D. For research, ask Claude to separate sourced facts from its own inferences and to flag gaps.

E. For analysis, ask Claude to keep its answer creative and open-ended.

Q8. A user's prompt asks Claude to "list all the risks" in a project plan. Claude returns 35 risks of very uneven importance. What is the BEST next prompt?

A. Ask Claude to remove any risk it is not certain about.

B. Ask Claude to add more risks so the list is comprehensive.

C. Ask Claude to shorten the list to 10 risks.

D. Ask Claude to rate each risk by likelihood and impact using defined scales, group related risks, and return the top 10 with mitigations.

Domain 2: Output Evaluation and Validation (21%)

Q9. Claude summarises a policy and states: "Employees may carry over up to 10 unused leave days." The policy says "up to 5 days, or 10 days with manager approval." How should the reviewer classify and handle this?

A. As a hallucination: discard the summary and write it manually.

B. As acceptable, because 10 days is possible under the policy.

C. As a style issue: add "in some cases" after the number.

D. As a material inaccuracy caused by dropping a condition: correct the statement to include the condition, then check other conditional rules in the summary.

Q10. Which TWO signals in a Claude response MOST warrant independent fact-checking before use? **(Select TWO)**

A. Use of headings and bullet points.

B. Precise statistics or dates that were not in the material provided to Claude.

C. References to named studies, laws, or sources that the user did not supply.

D. A suggestion that the user verify a figure.

E. A clear statement of which parts are based on the uploaded document.

Q11. A legal team uses Claude to summarise case law for an internal memo. What is the MINIMUM appropriate validation before the memo is relied upon?

A. Check that the memo is well organised and easy to read.

B. Check a random sample of two citations out of the twelve in the memo.

C. Ask Claude to confirm that all citations are real.

D. Check that every cited case exists and supports the stated proposition, with a qualified lawyer reviewing the conclusions.

Q12. Claude writes a competitor analysis stating "Competitor X has the largest market share." The analyst's own data shows Competitor X third by revenue but first by number of units sold. What is the MOST accurate assessment?

A. The claim is false and must be deleted.

B. The claim is accurate, since it is true by at least one measure.

C. The claim is ambiguous and potentially misleading. It should specify the metric (units) and be checked against the analyst's data.

D. The claim is a bias issue, because Claude favours well-known brands.

Q13. A Claude-generated customer-churn analysis concludes that "customers who contact support churn more, so support interactions cause churn." What is the MOST important flaw to flag?

A. The analysis should have used a chart instead of text.

B. The data should have been anonymised first.

C. It infers causation from correlation. Customers with problems may both contact support and churn, so the conclusion needs reframing or further analysis.

D. The conclusion is too short.

Q14. A team asks Claude to translate a product safety notice into Spanish for customers. What is the MOST appropriate validation step?

A. Publish it, since modern models translate reliably.

B. Ask Claude to rate its translation quality from 1 to 10.

C. Have a fluent speaker with product knowledge review the translation, particularly the safety-critical instructions, before publishing.

D. Translate it back into English with Claude and compare the result with the original.

Q15. Claude drafts a project-closure report saying "all 14 deliverables were completed on time." The project tracker shows 12 completed on time, 1 late, and 1 descoped. What does this MOST likely indicate, and what should be done?

A. The tracker is probably wrong, because Claude analysed the full context.
B. It is a rounding difference and acceptable in a closure report.
C. The report should simply remove the number of deliverables.
D. The summary overstated success by smoothing over exceptions. Correct it from the tracker and check whether exceptions were omitted elsewhere.

Q16. A manager has two Claude drafts of a client proposal: Draft 1 is persuasive but makes two commitments the firm cannot meet; Draft 2 is accurate but dry. What is the BEST approach?

A. Use Draft 2 as the factual base, then ask Claude to improve persuasiveness while keeping every commitment within what the firm can deliver, and review the final version.
B. Merge the two drafts paragraph by paragraph and send the result.
C. Send Draft 1, since persuasive proposals win more business.
D. Send Draft 2 without changes, since accuracy is all that matters.

Q17. A user needs to share results of a Claude-assisted data analysis with a team that will paste the figures into their own spreadsheets. Which format is BEST?

A. A bulleted list mixing figures and commentary.
B. A clean table or CSV-style structured output with labelled columns and units.
C. A narrative paragraph describing the figures.
D. A slide-style artifact with charts only.

Q18. When deciding whether a Claude output requires human review, which TWO factors are MOST relevant? **(Select TWO)**

A. Whether the output will reach an external audience or drive a decision.
B. Whether the output was produced quickly.
C. The consequences if the output is wrong (financial, legal, safety, or reputational).
D. The number of words in the output.
E. Whether the output contains formatting such as tables.

Q19. Claude writes an executive summary of a 60-page audit report. Which review approach BEST balances thoroughness and efficiency?

A. Re-read the full 60-page report and rewrite the summary manually.
B. Ask a colleague to read the summary and say whether it sounds right.
C. Trace every number and finding in the summary back to the report, and check that each high-severity finding in the report appears in the summary.
D. Read the summary for tone and grammar only, since the audit report has already been reviewed.

Q20. A Claude-drafted response to an angry customer is accurate but uses phrases such as "as previously stated" and "you failed to." What is the MOST appropriate refinement?

A. Remove all the facts so the message is purely apologetic.
B. Ask Claude to add more detail on why the customer is at fault.
C. Send it, since the facts are correct.
D. Ask Claude to keep the facts and resolution steps but rewrite them in an empathetic, non-blaming tone suited to the customer.

Q21. A user needs a reference list of 40 product SKUs with price and status for an internal team, and will update it weekly in conversation with Claude. Which output form is MOST suitable?

A. Project instructions containing the SKU list.
B. Bullet points in the chat, copied manually into an email each week.
C. A new inline paragraph each week.
D. A table maintained as an artifact that can be revised and exported.

Domain 3: Product and Model Selection (12%)

Q22. A team builds an internal tool that tags the language and sentiment of 2 million short app-store reviews each month. Which model family is MOST appropriate?

A. Haiku
B. Sonnet for half and Opus for the other half
C. The most capable available model, to maximise accuracy
D. Opus

Q23. A policy analyst must synthesise conflicting evidence from many long reports into a nuanced recommendation for ministers. Errors would be costly; turnaround is three days. Which choice is BEST?

A. Sonnet only, because it is the standard choice for business tasks.
B. Haiku, because faster responses allow more iterations.
C. Whichever model the analyst used last time, for consistency.
D. Opus, because the task requires deep reasoning over complex material and time is not tight.

Q24. Which statement BEST describes the trade-off across the Claude model families?

A. All three have identical capability and differ only in price.
B. Haiku is designed for long documents, and Opus is designed for short questions.
C. Haiku prioritises speed and cost efficiency, Opus prioritises maximum capability, and Sonnet balances the two.
D. Opus is always the right choice because quality matters in every task.

Q25. A user needs to produce an interactive, reusable calculator and a formatted policy document during a chat, and refine both over several turns. Which feature is MOST relevant?

A. Connectors
B. Artifacts
C. Project knowledge
D. Research mode

Q26. A team member pastes a 300-page technical manual into a chat, asks detailed questions over many turns, and later notices that answers miss details from early sections. Which TWO actions BEST manage this context limitation? **(Select TWO)**

A. Switch to the fastest model to process the manual more quickly.
B. Continue indefinitely in the same chat, since earlier content is always fully retained.
C. Add the manual to a Project's knowledge so relevant sections can be retrieved as needed, rather than relying on one long chat.
D. Paste the manual a second time in the same chat to reinforce it.
E. Ask targeted questions that name the section or topic, and verify answers against the manual.

Q27. A user regularly does unrelated tasks (trip planning, contract review, code snippets) in one long-running chat. What is the BEST practice recommendation?

A. Start a new chat (or use a dedicated Project) for each distinct task, so irrelevant earlier context does not interfere.

B. Use the most capable model, which can handle mixed context without issues.

C. Add a reminder at the top of every message saying which task is current.

D. Keep everything in one chat so Claude learns the user's preferences.

Q28. A team must choose between a standard chat and research mode for a question about "our internal travel policy's per-diem rates." The policy is in their Project knowledge. Which is MOST appropriate?

A. A standard chat outside the Project, asking Claude from general knowledge.

B. Research mode, because it is more thorough.

C. Research mode followed by a standard chat to compare answers.

D. A chat in the Project that holds the policy, since the answer comes from internal knowledge rather than the web.

Domain 4: Workflow Integration and Solution Design (16%)

Q29. A legal team wants to use Claude in its contract-intake process. Which design BEST integrates Claude while preserving accountability?

A. Claude approves low-value contracts automatically, and lawyers review only contracts over a set value.

B. Claude rewrites counterparty contracts directly into the company template and sends them back.

C. Lawyers use Claude only when they personally choose to, with no defined process.

D. Claude produces a first-pass issues list against the playbook for every contract, lawyers review and decide on each issue, and outcomes are logged to refine the playbook.

Q30. A COO wants to know whether Claude could cut the time spent preparing board packs. What is the BEST first step?

A. Buy additional licences so every executive can experiment.

B. Ask Claude to write the next board pack end to end and compare.

C. Survey other companies to see whether they use AI for board packs.

D. Map the current board-pack process, measure where time is spent, and identify which steps (for example drafting commentary or summarising reports) Claude could assist with.

Q31. Which TWO statements BEST communicate Claude's limitations to stakeholders in a balanced way? **(Select TWO)**

A. "Claude's limitations are the vendor's responsibility, not ours."

B. "Claude can produce confident but incorrect statements, so outputs that drive decisions need verification."

C. "Claude only sees the data we provide or connect it to; it doesn't automatically know our internal systems."

D. "Claude is always correct when it gives citations."

E. "Claude is never reliable enough for real work."

Q32. A customer-success team wants Claude to help draft quarterly business reviews (QBRs) for 120 accounts. What is the BEST solution design?

A. A standard QBR template and instructions in a Project; account data pulled from approved sources; Claude drafts each QBR; account managers verify figures and add relationship context before sending.

B. Claude drafts generic QBRs without account data, and account managers fill in the details.

C. Each account manager writes their own prompts and formats as they see fit.

D. Claude drafts all 120 QBRs from last quarter's QBRs and sends them automatically.

Q33. A team pilots Claude for drafting internal policy FAQs. After four weeks, drafting time fell by 60%, but reviewers found errors in 15% of answers, mostly from outdated source policies. What is the BEST next step?

A. Fix the root cause by cleaning up and versioning the source policies Claude uses, then re-measure accuracy before expanding.

B. Switch to a more capable model to eliminate the errors.

C. Expand to all departments now, since the time savings are significant.

D. Stop the pilot, because a 15% error rate is unacceptable.

Q34. An operations manager wants to use Claude to help plan a warehouse relocation. Which use is MOST appropriate?

A. Having Claude negotiate the new lease with the landlord by email.

B. Having Claude decide the relocation date and notify staff.

C. Generating a draft project plan with phases, dependencies, and risks from the manager's inputs, for the team to validate and adjust.

D. Using Claude only to proofread the final plan.

Q35. A marketing team wants to redesign its campaign-brief process with Claude. Currently, briefs arrive incomplete and creative teams lose days chasing details. Which design BEST addresses the root problem?

A. Have Claude fill in missing brief details with reasonable assumptions so work can start.

B. Have Claude write creative concepts from the incomplete briefs to save time.

C. Add a longer brief template and keep the rest of the process unchanged.

D. Use Claude at intake to check each brief against a required-fields checklist, flag gaps, and generate clarifying questions before the brief is accepted.

Q36. Which approach BEST describes using Claude to support solution development and iteration for a new internal process?

A. Ask Claude for the final design once, then roll it out to everyone.

B. Draft an initial design with Claude, test it with a small user group, feed their feedback back to Claude for targeted revisions, and repeat until the success criteria are met.

C. Design the process manually and use Claude only for the announcement email.

D. Ask Claude for ten different designs and launch all of them at once.

Q37. A finance team wants Claude to help reconcile vendor statements against the ledger. Which design is MOST appropriate?

A. Claude posts adjusting journal entries automatically for any differences it finds.

B. The team keeps reconciling manually, since financial data should never be used with AI.

C. Claude estimates the likely reasons for differences without seeing the ledger data.

D. Claude compares the exported statement and ledger data, lists matched and unmatched items with reasons, and an accountant investigates and resolves the exceptions.

Q38. Which TWO indicators would BEST show that a Claude-augmented workflow should be scaled to more teams? **(Select TWO)**

A. A documented process (inputs, configuration, review steps) that other teams can follow.

B. The pilot ran without any review step, which proves Claude is reliable.

C. Claude produced more words of output than staff previously did.

D. Pilot users report that Claude feels impressive.

E. Measured time or cost savings in the pilot, with quality held steady or improved.

Domain 5: Configuration and Knowledge Management (12%)

Q39. A customer-support lead sets up a Project to help agents answer billing questions. Which configuration is BEST?

A. Instructions that define the role, the tone, and the response structure, tell Claude to answer only from the attached billing policies, say when information is not covered, and escalate refunds above a threshold; current billing policies in knowledge.

B. All historical support tickets (50,000) uploaded as knowledge, with no instructions.

C. Instructions that say "Answer billing questions helpfully," with no knowledge files, so Claude uses general knowledge.

D. Instructions that tell Claude to approve refunds when customers are unhappy.

Q40. Project instructions currently say: "Use a friendly tone. Use a formal tone for executives. Always be brief. Provide full detail on risks." Agents get inconsistent outputs. What is the BEST improvement?

A. Add more instructions to cover additional scenarios.

B. Rewrite the instructions to resolve the conflicts with clear rules and precedence (for example: "Default friendly; formal when the audience is executive. Be brief, except always list risks in full.")

C. Move the instructions into a knowledge file.

D. Delete the instructions and let each agent prompt as they wish.

Q41. Which TWO items are BEST suited to Project knowledge rather than Project instructions? **(Select TWO)**

A. The current product catalogue and pricing sheet.

B. "Escalate any legal question to the legal team."

C. The company's 25-page brand style guide.

D. "Format every answer with a summary line, then bullet points."

E. "Always respond in British English."

Q42. A Project used by 30 sales staff has knowledge files uploaded by many people over two years. Answers increasingly cite discontinued products. What is the BEST corrective action?

A. Have an owner audit the knowledge files, remove or archive outdated material, standardise file naming with dates and versions, and set a periodic review schedule.

B. Create a new Project and copy all existing files into it.

C. Upload the current catalogue again, so there are more current documents than old ones.

D. Add an instruction asking Claude to ignore discontinued products.

Q43. An employee wants to connect their Google Drive so Claude can help draft documents from their team's files. Which consideration is MOST important?

A. Disable Drive permissions temporarily while using Claude.

B. Make sure every folder in the Drive is shared with everyone so Claude can see them.

C. Confirm the connector is approved by the organisation, and understand that Claude will access files according to the employee's existing Drive permissions.

D. Copy all files into a personal folder first, to simplify access.

Q44. A team lead notices that Claude in their Project often answers questions outside its intended scope (for example personal tax questions), which creates risk. What is the BEST configuration change?

A. Remove all knowledge files so Claude has less to talk about.

B. Update the instructions to define the Project's scope and tell Claude how to respond to out-of-scope requests (for example, decline politely and point to the appropriate resource).

C. Remind users at team meetings not to ask off-topic questions.

D. Switch the Project to a smaller model, which answers fewer questions.

Q45. A Project's instructions were revised last week, and since then output quality has dropped. What is the BEST way to manage this?

A. Delete the Project and recreate it from memory.

B. Compare the outputs before and after the change on the same test prompts, restore or adjust the instruction text, and keep a version history for future changes.

C. Ask users to write longer prompts to compensate.

D. Keep the new instructions, since recent changes are usually improvements.

Domain 6: Governance, Risk, and Responsible Use (15%)

Q46. A recruiter wants to use Claude to rank 200 applicants and automatically reject the bottom 150. What is the MOST appropriate assessment?

A. Acceptable if Claude is told to be fair.

B. Acceptable if the recruiter spot-checks 10 rejected applicants.

C. High risk. Automated rejection affects people's opportunities, may embed bias, and may be regulated, so humans should make the decisions using job-relevant criteria, with Claude assisting at most.

D. Acceptable because it treats every applicant consistently.

Q47. Which TWO actions BEST reflect data minimisation when using Claude to draft a client letter about a billing dispute? **(Select TWO)**

A. Paste the client's full CRM profile so Claude has complete context.

B. Paste the full payment-card number so Claude can reference it accurately.

C. Include the client's prior complaints about unrelated matters, for tone guidance.

D. Include only the details needed for the letter (the invoice number, the disputed amount, and the timeline).

E. Omit unrelated personal data such as the client's date of birth and full payment-card number.

Q48. Which use case is MOST clearly inappropriate under typical organisational AI policies?

A. Using Claude to draft a project-update email that the sender reviews.

B. Using Claude to brainstorm team-building activities.

C. Using Claude to write messages that impersonate a specific real colleague, to obtain information from another department.

D. Using Claude to summarise a public research paper.

Q49. A healthcare administrator wants to upload patient records to Claude to draft referral letters. Which consideration is MOST important?

A. Ask patients afterwards whether they mind.

B. Confirm the tool and configuration are approved for health data under applicable regulations and organisational policy, and use only the minimum necessary information.

C. Check that Claude writes in a professional medical tone.

D. Upload all records at once to save time.

Q50. A teacher uses Claude to generate feedback on student essays. Which practice BEST reflects responsible use?

A. Let Claude assign final grades so that grading is objective.

B. Use Claude's feedback as a draft, review and adjust it for each student, follow school policy on AI use and student data, and avoid uploading unnecessary identifying information.

C. Send Claude's feedback directly to students, since it is consistent.

D. Upload students' full names and records so feedback is personalised.

Q51. Your organisation's AI policy prohibits entering confidential M&A information into any AI tool. A director asks you to use Claude to summarise a confidential acquisition target's financials "just this once." What is the BEST response?

A. Comply, and delete the chat afterwards.

B. Decline, explain the policy, and suggest raising an exception request through the governance process or using an approved alternative.

C. Comply, since the director outranks the policy owner.

D. Comply, but strip out the company name first.

Q52. Which statement BEST describes an ethical concern specific to relying on AI-generated content?

A. Users may over-trust fluent, confident output and skip verification (automation bias), which allows errors to spread.

B. AI-generated content cannot be edited after it is produced.

C. AI-generated content always uses more words than human writing.

D. AI-generated content is always less accurate than any human's work.

Q53. Which TWO practices BEST align with following organisational AI governance standards? **(Select TWO)**

A. Following the required review and disclosure steps for the type of output.
B. Deciding personally which data classifications are safe to share.
C. Skipping review for outputs that look correct.
D. Using only approved tools and features for work data.
E. Using personal accounts when the approved tool is slower.

Q54. A public-sector team wants to use Claude to draft responses to citizen inquiries. Which safeguard is MOST important?

A. Responses are kept short to reduce the chance of errors.
B. Responses are sent automatically to cut backlogs.
C. Claude is told to sound official.
D. Staff review responses for accuracy and fairness before sending, using approved tooling, with escalation rules for complex or sensitive cases.

Domain 7: Troubleshooting and Optimization (10%)

Q55. Claude's answers in a Project are correct but each one begins with a long preamble that users skip. Which fix is MOST efficient?

A. Remove the knowledge files.
B. Switch to a different model.
C. Add a Project instruction to lead with the direct answer, then supporting detail, with an example of the preferred structure.
D. Ask each user to say "skip the intro" in every message.

Q56. A user complains that Claude "didn't follow the brief." On inspection, the prompt was a forwarded email chain with the actual request buried in the third message. What is the BEST resolution?

A. Rewrite the prompt with the request stated up front and clearly, and include the email chain as labelled background.
B. Ask Claude to read more carefully.
C. Report the issue as a model malfunction.
D. Use a larger model to parse email chains better.

Q57. Which TWO steps reflect a sound process for diagnosing an underperforming prompt? **(Select TWO)**

A. Collect several failed outputs and identify the common pattern in the failures.
B. Change one element of the prompt at a time and compare results on the same test inputs.
C. Judge success by a single new output.
D. Rewrite the entire prompt and the configuration at once.
E. Assume the model is at fault before checking the inputs.

Q58. A team's workflow has Claude draft a report, a reviewer fix formatting, Claude summarise it, and a second reviewer re-fix formatting. Which optimisation is MOST effective?

A. Remove all review steps, since formatting is minor.

B. Ask Claude to produce the report in plain text so there is nothing to format.

C. Specify the required format in the instructions (with a template) so drafts are produced correctly the first time, and keep a single substantive review.

D. Add a third reviewer to catch remaining formatting issues.

Q59. After users gave feedback that summaries were "too generic," a team added more instructions, but the summaries are now overly long. What is the BEST adjustment?

A. Ask users to stop reading the summaries in full.

B. Identify the specific elements users valued (for example decisions, figures, and owners), require those, and set an explicit length limit.

C. Remove all instructions and start over.

D. Ignore the feedback, since you cannot satisfy everyone.

Q60. A team uses Opus for all tasks, including simple email triage, and costs are rising. Quality on complex analysis must be maintained. What is the BEST optimisation?

A. Keep using Opus for everything, since quality always justifies cost.

B. Switch everything to Haiku immediately.

C. Route simple, high-volume tasks such as triage to Haiku or Sonnet, and keep Opus for complex analysis where it adds measurable value.

D. Reduce the number of users who can access Claude.

Answer Grid

Q1	**B**	Q2	**C**	Q3	**AE**	Q4	**B**	Q5	**D**	Q6	**D**
Q7	**BD**	Q8	**D**	Q9	**D**	Q10	**BC**	Q11	**D**	Q12	**C**
Q13	**C**	Q14	**C**	Q15	**D**	Q16	**A**	Q17	**B**	Q18	**AC**
Q19	**C**	Q20	**D**	Q21	**D**	Q22	**A**	Q23	**D**	Q24	**C**
Q25	**B**	Q26	**CE**	Q27	**A**	Q28	**D**	Q29	**D**	Q30	**D**
Q31	**BC**	Q32	**A**	Q33	**A**	Q34	**C**	Q35	**D**	Q36	**B**
Q37	**D**	Q38	**AE**	Q39	**A**	Q40	**B**	Q41	**AC**	Q42	**A**
Q43	**C**	Q44	**B**	Q45	**B**	Q46	**C**	Q47	**DE**	Q48	**C**
Q49	**B**	Q50	**B**	Q51	**B**	Q52	**A**	Q53	**AD**	Q54	**D**
Q55	**C**	Q56	**A**	Q57	**AB**	Q58	**C**	Q59	**B**	Q60	**C**

Answer Key with Explanations

Domain 1: Prompting and Task Execution

Q1. Answer: B
Analysis tasks need a defined purpose, evaluation criteria, and the shape of the deliverable (B). C pushes Claude toward exactly the exhaustive description the officer got. D helps with tone, but without criteria Claude still cannot weigh the quotes. A changes the format but still provides no basis for a decision.

Q2. Answer: C
Supplying all the constraints up front, clearly separated, produces a compliant first draft (C). A leaves out the approved messages and the legal constraints. B builds in a rework loop and risks missing prohibited phrases. D may carry over wording that does not fit the current situation.

Q3. Answer: A, E
Explicit boundaries on what to keep (E) and a concrete definition of "improve" (A) stop scope creep into the content. C is vague. D increases the risk of changed facts. B helps review after the fact, but it does not prevent the problem.

Q4. Answer: B
Specific context, policy constraints, tone, and anticipated responses produce a script that can actually be used (B). C is a reasonable fallback, but it leaves the hardest tailoring to the user. D is valuable later, once the content exists. A defeats the purpose of a PIP conversation.

Q5. Answer: D
Staged decomposition with checkpoints improves depth and structure (D). A adds words, not depth. B can help with reasoning tasks, but it does not create review points. C produces a draft stitched together from parts that may not fit.

Q6. Answer: D
Examples, plus a short explanation of each section, are the most reliable way to reproduce a specific house format (D). C leaves too much to inference. A produces a generic format instead of the team's own. B adds a conversion step without guaranteeing the format.

Q7. Answer: B, D
Research benefits from distinguishing sources from inference (D), and drafting benefits from audience and purpose constraints (B). C ignores task type. A applies a research constraint to creative work. E applies a brainstorming posture to analysis.

Q8. Answer: D
Adding a prioritisation method and a clear output shape turns a raw list into decision-ready analysis (D). C cuts the list with no basis for choosing which risks to keep. B makes the noise worse. A filters on Claude's own confidence instead of on risk significance.

Domain 2: Output Evaluation and Validation

Q9. Answer: D
Omitting the condition changes the meaning of the rule. Correct it and look for the same pattern elsewhere (D). B accepts a misleading statement. A overreacts, since most of the summary may be fine. C remains vague and still misleading.

Q10. Answer: B, C
Specific facts (B) and named sources (C) that did not come from supplied material are where fabrication risk is highest. A is formatting. E and D are signs of transparency, not problems.

Q11. Answer: D
Legal citations are high-risk for fabrication and must each be verified, with qualified review of the conclusions (D). A checks presentation only. C is self-verification. B leaves ten citations unchecked in a high-stakes document.

Q12. Answer: C
"Market share" without a stated metric can mislead. Precision and alignment with the data fix it (C). A discards a claim that is true when properly qualified. B ignores the reader's likely interpretation. D misdiagnoses the issue.

Q13. Answer: C
Evaluating reasoning quality includes spotting causal claims that the evidence does not support (C). A and D are presentation issues. B may matter for privacy, but it is not the analytical flaw being asked about.

Q14. Answer: C
Safety-critical, customer-facing content needs qualified human review (C). D is a useful supplementary check, but it can miss errors in tone, terminology, or regional usage. A skips validation. B is self-assessment.

Q15. Answer: D
Summaries can drift toward a tidy, positive narrative and gloss over exceptions. The source of truth wins, and you should check for similar omissions (D). A reverses the source of truth. B accepts a misstatement. C hides the issue.

Q16. Answer: A
Comparing and refining outputs means keeping accuracy as the constraint and then improving the style (A). C creates commitments the firm cannot meet. D misses an easy quality gain. B may import the inaccurate commitments.

Q17. Answer: B
Data that will be reused needs a structured format with clear labels and units (B). C, D, and A require manual extraction and invite transcription errors.

Q18. Answer: A, C
Review intensity should scale with the impact of errors (C) and with exposure or decision use (A). D, B, and E are not risk indicators.

Q19. Answer: C
Targeted verification of accuracy (numbers and findings) and completeness (high-severity items) is efficient and thorough (C). D misses errors introduced in summarising. A defeats the efficiency gain. B gives an impression, not verification.

Q20. Answer: D
Refining for the audience means keeping the substance while adjusting the tone (D). C risks escalating the complaint. A removes the needed information. B makes the tone problem worse.

Q21. Answer: D
A structured, standalone, iteratively maintained deliverable fits a table in an artifact (D). C and B are hard to maintain. A misuses instructions, which govern behaviour and are not the place for data deliverables.

Domain 3: Product and Model Selection

Q22. Answer: A
Simple, very high-volume tagging favours the fastest, most cost-efficient family (A). D and C are expensive and slower, with little benefit on a simple task. B adds cost and inconsistency.

Q23. Answer: D
Complex, high-stakes reasoning with time available favours the most capable family (D). B optimises the wrong variable. A is a sensible general default, but this task's priorities point higher. C is not a selection criterion.

Q24. Answer: C
C is the accepted characterisation. A is false. D ignores cost and speed. B reverses or misstates the families' purposes.

Q25. Answer: B
Artifacts hold substantial, standalone outputs, such as documents and simple interactive content, that can be iterated on (B). Research mode (D) gathers information. Project knowledge (C) stores reference material. Connectors (A) link Claude to external data.

Q26. Answer: C, E
Persisting large reference material in project knowledge (C) and asking targeted questions with verification (E) work within the context limits. D consumes more context. A does not address the limit. B is false.

Q27. Answer: A
Separating unrelated work keeps context clean and relevant (A). D risks contamination between tasks and context limits. C is a weak workaround. B does not remove the interference from irrelevant context.

Q28. Answer: D
Internal-policy questions should be grounded in the internal source (D). Research mode (B, C) targets external sources and may surface irrelevant public material. A has no access to the policy.

Domain 4: Workflow Integration and Solution Design

Q29. Answer: D
A defined AI step, human decision, and feedback loop is sound integration (D). A delegates approval to Claude. C lacks consistency and governance. B sends unreviewed external communications and changes.

Q30. Answer: D
Analysing the current workflow and pinpointing where Claude fits comes before tooling or pilots (D). A invests before there is a use case. B is a high-risk pilot on a sensitive deliverable. C gathers secondary information but not the analysis this process needs.

Q31. Answer: B, C
Accurate limitations are possible errors that require verification (B) and dependence on supplied or connected data (C). E understates Claude's value. D is false. A misplaces accountability.

Q32. Answer: A
Standardised configuration, grounded data, and human verification scale well (A). D uses stale data and sends without review. C produces inconsistency. B delivers little value because the account-specific content is missing.

Q33. Answer: A
The errors trace to stale inputs, so fixing the inputs and re-measuring is the evidence-based iteration (A). C scales a known defect. D gives up on a fixable problem. B misdiagnoses it, since a better model cannot correct outdated source policies.

Q34. Answer: C
Claude supports planning by structuring and drafting from inputs, with humans validating the result (C). B and A hand Claude consequential decisions or external actions. D under-uses its planning support.

Q35. Answer: D
The problem is incomplete inputs. A Claude-assisted completeness check at intake fixes it at the source (D). A bakes in unverified assumptions. B builds on incomplete information. C may help, but it does not enforce completeness.

Q36. Answer: B
Iterative build-test-refine cycles with clear criteria are the core of solution design (B). A skips validation. D creates confusion. C misses design support.

Q37. Answer: D
Claude speeds up matching and exception reporting, and the accountable human resolves the exceptions (D). A removes control over financial postings. C is not grounded in the data. B is overly restrictive where approved tools and controls exist.

Q38. Answer: A, E
Evidence of value (E) and a repeatable, documented process (A) justify scaling. D is anecdotal. C measures volume, not value. B is a risk, not a success indicator.

Domain 5: Configuration and Knowledge Management

Q39. Answer: A
Clear instructions grounded in curated, current knowledge, with explicit boundaries, form an effective configuration (A). C invites generic or invented answers. B is noisy, unstructured knowledge with no guidance. D gives Claude a decision it should not make.

Q40. Answer: B
Conflicting or ambiguous instructions cause inconsistency. Explicit precedence resolves it (B). D loses standardisation. A may add further conflicts. C reduces the instructions' directive role.

Q41. Answer: A, C
Reference material such as catalogues (A) and style guides (C) belongs in knowledge. Short behavioural rules (E, D, B) belong in instructions.

Q42. Answer: A
Stale knowledge needs curation, naming conventions, and ongoing ownership (A). D relies on Claude knowing which products are discontinued. C leaves the conflicting files in place. B moves the problem into a new Project.

Q43. Answer: C
Use approved connectors and understand that they respect existing permissions (C). B, D, and A weaken data-access controls.

Q44. Answer: B
Scope boundaries and fallback behaviour belong in the instructions (B). A breaks the in-scope use. C relies on people's memory. D does not establish a scope.

Q45. Answer: B
Treat instructions as versioned configuration that you test and can roll back (B). D ignores the evidence. A loses working configuration. C shifts the burden onto users without fixing the cause.

Domain 6: Governance, Risk, and Responsible Use

Q46. Answer: C
Consequential decisions about people need human accountability and bias safeguards, and some jurisdictions regulate automated hiring decisions (C). B provides minimal oversight. A is an instruction, not a safeguard. D confuses consistency with fairness.

Q47. Answer: D, E
Data minimisation means providing only what the task requires (D) and excluding unnecessary sensitive data (E). A, C, and B add unnecessary personal or sensitive data.

Q48. Answer: C
Impersonation to obtain information is deceptive and a social-engineering risk (C). The others are routine, reviewable uses.

Q49. Answer: B
Health data carries specific regulatory obligations. Approved tooling and minimum-necessary use come first (B). C is about quality, not compliance. D over-shares. A is retroactive and inadequate.

Q50. Answer: B
Human review, policy compliance, and data minimisation reflect responsible use (B). C and A remove educator accountability. D over-shares personal data.

Q51. Answer: B
Policy applies regardless of seniority. Use the formal exception or approval route (B). C ignores governance. D partial anonymisation may still expose confidential deal information. A does not undo the disclosure.

Q52. Answer: A
Automation bias, meaning over-reliance on confident AI output, is a well-recognised ethical and operational risk (A). C and B are false. D is an overgeneralisation.

Q53. Answer: A, D
Approved tools (D) and mandated review and disclosure steps (A) are core governance practices. E, B, and C bypass governance.

Q54. Answer: D
Citizen-facing public communications need accountable review, approved tools, and escalation paths (D). B removes oversight. C and A are style choices, not safeguards.

Domain 7: Troubleshooting and Optimization

Q55. Answer: C
A recurring format issue is fixed once in the configuration (C). D repeats the fix in every message. B and A do not address the cause.

Q56. Answer: A
The root cause is an unclear input structure. A clear request plus labelled context fixes it (A). C and D misdiagnose the cause. B is vague.

Q57. Answer: A, B
Pattern analysis (A) and controlled, one-variable-at-a-time testing (B) isolate the cause. D makes it impossible to tell which change mattered. C is anecdotal. E skips the most common cause.

Q58. Answer: C
Preventing the defect upstream removes the repeated rework while keeping meaningful review (C). D adds cost. A removes needed quality control. B changes the requirement instead of meeting it.

Q59. Answer: B
Adjusting from feedback means targeting what users valued while constraining the side effect (B). C discards progress. D and A avoid improving the summaries.

Q60. Answer: C
Matching models to tasks optimises cost while protecting quality where it matters (C). B risks degrading complex analysis. A ignores the cost problem. D reduces value instead of optimising it.