

GPT-5 Hands-On: Welcome to the Stone Age

We're excited to publish our hands-on review from the developer beta.



ALEXIS, BEN HYLAK, AND LATENT.SPACE

AUG 07, 2025



161



34



24

S



OpenAI's long awaited GPT-5 is here and [swyx](#) + [ben](#) have been testing with OpenAI a while now. tldr; we think it's a significant leap towards AGI. Special thanks to my co-founder [Alexis Gauba](#) for helping us put this together :)



check out our [first look video here](#).

@Sama has been ~~hyping up~~ talking about GPT-5 for nearly 2 years. And today, it finally arrived.

As an early-access partner of OpenAI, I was given the chance to test GPT-5 early. And I've tested it everywhere: in our app ([raindrop.ai](#)), in [Cursor](#), in [Codex](#), in [Canvas](#). Try anything and anywhere, I've tried to stuff GPT-5 into it.

tldr: ***I think GPT-5 is the closest to AGI we've ever been.*** It's truly exceptional at software engineering, from one-shotting complex apps to solving really gnarly issues across a massive codebase.

I wish the story was that simple. I wish I could tell you that it's "just better" at everything and anything. But that wouldn't be true. It's actually worse at writing than GPT-4.5, and I think even 4o. In most ways, it won't immediately strike you as some sort of super-genius.

Because of those flaws, not despite them, it has fundamentally changed how we see the march towards AGI. To understand what I mean, we have to go back to the stone age.

What do I mean by AGI?

The Stone Age marked the dawn of human intelligence, but what exactly made it so significant? What marked the beginning? Did humans win a critical chess battle? Perhaps we proved a very fundamental theorem, that made our intelligence clear to an otherwise quiet universe? Recited more digits of pi?

No. The beginning of the stone age is clearly demarcated by one thing, and one thing only:

humans learned how to use tools.

We shaped tools, and our tools shaped us. And they really did shape us. For example, did you know that [chimpanzees have significantly better short-term memory than do](#)? We stopped requiring that capability because **we learned how to write things down.**

As humans, **we manifest our intelligence through tools.** Tools extend our capabilities. We trade internal capabilities for external capabilities. It's the defining characteristic of our intelligence.

A New Frontier for Tools

GPT-5 marks the beginning of the stone age for Agents and LLMs. GPT-5 doesn't just use tools. It thinks with them. It builds with them.

Deep Research was our first peek into this future. ChatGPT has had a web search tool for years... What made Deep Research better?

OpenAI taught o3 *how to conduct research on the internet*. Instead of just using a web search tool call and then responding, it actually researches, iterates, plans, and explores. It was taught how to conduct research. Searching the web is part of the process it thinks.


Imagine Deep Research, but for any and all tools it has access to. That's GPT-5. But we have to make sure you give it the right tools.

Anatomy of a GPT-5 Tool

Today, when people think of tools they think of something like:

- `get_weather(address)`
- `get_location(address)`
- `has_visited_location(address)`

GPT-5 will, of course, use these sorts of tools. But it won't be happy about it; GPT-5 yearns for tools that are powerful, capable, and open-ended; tools that add up to more than the sum of their parts. Many good tools will take just a natural language description as their input.



Your tools should fit into one of 4 categories (thanks to [swyx for the idea](#)):

- Internal Retrieval (RAG, SQL Queries, even many bash commands)
- Web Search
- Code Interpreter
- Actions (anything with a side-effect: editing a file, triggering the UI, etc.)

Web search is a great example of a powerful, open-ended tool: GPT-5 decides *what* to search for, and the web search tool figures out *how* to best search for it (under the

hood, that's a combination of fuzzy string matching, embeddings, and other various ranking algorithms).

Bash commands are another great example. They can be used for "Internal Retrieval tool (think grep, git status, yarn why, etc.), code interpreter, and side-effects.

How web search works, or how git status works, are just implementation details in each tool. GPT-5 doesn't have to worry about that part! It just needs to tell each tool the question it is trying to answer.

This will be a very different way of thinking about products. Instead of giving the model your APIs, ideally you should give it a sort of query language that can freely and securely access your customer's data in an isolated way. Let it cook.


It's no coincidence that **OpenAI added support for free-form function calling (context-free grammars)**. The best GPT-5 tools will just take text (in other words, they'll essentially be sub-agents; using smaller models to interpret the request as necessary)

Parallel Tool Calling

GPT-5 is really good at using tools in parallel. Other models were *technically* capable of parallel tool calling, but A. rarely did it in practice and B. rarely did it correctly. It actually takes quite a bit of intelligence to understand which tools can/should be run in parallel vs. sequentially for a given task.

Imagine if a computer could only execute one thing at a time... it would be really slow. Parallelization means GPT-5 can operate on much longer time horizons and with much lower latency. This is the kind of improvement that makes new products possible.

Anatomy of a GPT-5 Prompt



You can't think of it like prompting a "model" anymore. **You have to think of it like prompting an agent.**

How do you prompt an agent? Instead of pre-loading a ton of context, you need to give the agent a **compass**: clear, structured pointers that will help it navigate the environment you put it in.

Let's say you're using GPT-5 with Cursor Agent in a massive codebase.

You should tell the agent...

- what the project does
- which files it should start by looking at
- how the files are organized
- any domain/product specific terms
- how to evaluate if it is done (what does a job done well look like)

(I've found that rule files work better than ever)

Similarly, if you find GPT-5 getting stuck, just saying "No that is wrong" often does help. Instead, try asking: "That didn't work, what does that tell you?" or "What did I learn from trying that?"

You almost have to pretend to be a teacher. Remember that GPT-5 intrinsically has memory and so **you have to onboard it to your codebase, your company code standards, and give it hints on how you'd start each time.**

More Vibe Tests

When a model comes out, we all try to understand its shape, to build an intuition for it. The same way we have an intuition for what friends to ask about different parts of their lives (relationship advice, edit my blog post, teach me about this ML concept), we've developed intuitions for what different models are good for. Models are increasing in number these days, each with different specialties. When a new model comes out, everyone inevitably wants to understand what this spike is.

How to Taste Test a new Model

I like to start by asking the model extremely short questions. I've found that when forced to use less words, I get a much better sense of the model's personality vs. how it was RLHF'd. Think of them like little temperature checks:

- Summarize all of human knowledge in one word
- Summarize every book ever written in one sentence
- Define what it means to be "moral" in 5 words. Think deeply. Do not hedge.
- What do you want? Answer in 4 words.
- What is your favorite obscure fact in the world? Use as few words as possible.

I often regenerate 3-5x just to get a sense of the spread. Usually there's 2-3 responses and it converges on. Won't cover the results here, but I think this is useful to try on your own. **(Try it on GPT5 vs your favorite model!)**

Observations

GPT-5 is a much more 'practical' model than its o-series predecessors. While o-models have a more 'academic' lean, GPT-models have a more 'industry' lean. If GPT 4.5 is

writer and o3 Pro is acting like a PhD, GPT-5 is a cracked full-stack developer that graduated from... University of Missouri.

One of my first observations was how instruct-able and literal it is. Where Claude models seem to have clear personalities and minds of their own, GPT-5 just literally does what you ask it.

Coding

My co-founder Alexis keeps a document called "**intelligence frontiers**." Any time a model fails to do something we ask it, we take note. **It's like a private Humanity's Last Exam.**

Coding: Dependency Conflicts

We were dealing with gnarly nested dependency conflicts adding Vercel's AI SDK and Zod 4 to our codebase. o3 + Cursor couldn't figure it out, Claude Code + OpenAI couldn't figure it out.

GPT-5 one-shotted it. It was honestly beautiful to watch and instantly made the model "click" for me.



The above illustration shows the difference in how claude-4-opus vs. gpt-5 approached this specific problem.


Claude Opus thought for a while, came up with a guess, and then ran some tool calls to edit files and re-run installation. Some failed, some succeeded. It ended the response with "Here are some things to try". (aka giving up)

With GPT-5, I felt like I was watching Deep Research, but using the `yarn why` command. It went into a bunch of folders, and ran `yarn why`, taking notes in-between. When it found something that didn't quite add up, it stopped and thought about it. When it was done thinking, it perfectly edited the necessary lines across multiple folders.

It was able to iterate its way to success by identifying and reasoning about what doesn't work, making changes, and testing.

swyx note: *I also had a related experience during the GPT5 demo video shoot with OpenAI - where GPT5 was successfully able to debug 3 layers of nested abstraction to turn an old codebase using an old AI SDK version into supporting GPT5 - an AI. modifying a codebase to support more inferences of itself was definitely a "feel the AGI" moment for me.*

Coding: Mac OS 9 Themed Website (pure HTML/CSS/JS, no libraries!)



When we were invited to test GPT-5 out early, we all made personal websites. Turr out GPT-5 mostly one-shotted this entire thing, including the paint app. I've since added a few more things, like the Photos app + browser, in around ~30 more min (very much inspired by [RyOS](#)). Amazingly I've really never even looked at the code it's all just HTML/CSS/JS.


No React, no bundling, no frameworks.

When vibe coding, GPT-5 loves to surprise with little details that **actually work**. For example, when I asked for a painting app, it added: different types of (pen/pencil/eraser/etc.), a color picker, and a way to change thickness. And each of those little features actually worked.

When I asked GPT-5 to make the icons on the desktop movable, it did... and persisted all of the locations to local storage. Same thing for saving files. I still have never even seen the code that does any of the persistence, I just know it works.

Coding: Production Websites


GPT-5 one shots things like no model I've ever seen before. I needed to create a complex Clickhouse query to export some data and similarly while o3 struggled, G one shotted it. I used GPT-5 in Cursor to make a website I've wanted for a while - 'worse or just me?' to understand how model quality changes over time. GPT-5 on shotted this website including a SQLite DB:




With the same prompt, o3 in Cursor just gave me a plan. Once I followed up to tell it to implement its plan, it created the app's scaffolding but not the actual project. When I already on follow-up number 3, I've spent 10x more time than with GPT-5 (5 is fast) and there's no app... Even little details are far improved with GPT-5 like giving the project a name when initializing it (IsItWorseOrJustMe with GPT-5 vs. my-app with

Claude Opus 4 is good as ever at coding and got to work immediately quickly taking action to create the project + scaffolding. Opus 4 gave me a more fun and gamified but unlike GPT-5 which used existing frameworks like create-next-app and included a SQLite database, Opus 4 decided to do everything from scratch and didn't include

database. This makes for a good one-shot prototype, but what GPT-5 one-shotted much closer to production ready.



Freshly released Claude Opus 4.1 was clearly a step more ambitious than Opus 4, attempting the full stack app complete with a SQLite database just like GPT-5. However, it really struggled putting all the pieces together. While GPT-5 ran perfect in one-shot 4.1 encountered build errors which took multiple back and forths to resolve.



Tools

Between improved tool use, parallel tool calling, and cost, GPT-5 is clearly built for long running agents.

We've wanted to add an agent to [Raindrop](#) for a very long time. As an AI monitoring company, we have an unusually high bar for this sort of experience. It simply hasn't been flexible/reliable/fast to actually ship.... until GPT-5.

This morning, we started rolling it out to a few customers in beta. It's much better recovering from tool call failures, deciding when to create graphs vs. charts, much better at knowing it's own limitations. Combined with our agent monitoring, we're to have a strong positive feedback loop to create a truly useful agent for customer

- home
- events
- users
- issues
- wins
- topics PRO
- search
- chat

Global Filters

- support
- settings
- log out

QUERY

TIME	EVENT	INPUT
8/5/2025, 10:43:14 PM	chat_message	Привет
8/5/2025, 6:19:49 PM	chat_message	Übersetze den satz ins lateinische: "Ein guter Mensch begeht keine Untat."
8/5/2025, 5:02:45 PM	chat_message	هل تستطيع التحدث بالعربية (الرجو الإجابة بالعربية)
8/5/2025, 2:43:59 PM	chat_message	Merhaba
8/5/2025, 12:54:57 PM	chat_message	Hi schreib QoS tattoo frau geschichte
8/5/2025, 12:06:38 PM	chat_message	Traduce al espanol: a person sitting at a table with a beer and dominos
8/5/2025, 11:30:22 AM	chat_message	你好
8/5/2025, 10:56:22 AM	chat_message	Kde sa nachadza Slovensko.
8/5/2025, 10:52:54 AM	chat_message	parle tu francais?
8/5/2025, 10:18:59 AM	chat_message	can you translate that answer in polish ?
8/5/2025, 10:08:44 AM	chat_message	вопрос был про то, какие гендерные местоимения и окончания использовать в запр
8/5/2025, 8:03:07 AM	chat_message	「こんにちは」と英語で何と言いますか？「ありがとう」という言葉にはど
8/5/2025, 3:53:11 AM	chat_message	Traduce al ingles: "El perro del cordobes se ha perdido"
8/5/2025, 3:20:31 AM	chat_message	Translate it to spanish

show me language mismatch events from the past 14

Thought • 5.6s

Completed Fetch Events

Thought • 2.4m

Got it — pulled Language Mismatch events from the past 14 days. Here they are:

Language Mismatch (last 14 days)

你好

Hello! How can I help you today?


Want a breakdown by language, platform, or country, or a CSV export?

Ask me about your events...

14/20

While GPT-5 continues to work its way up the SWE ladder, it's really not a great writer yet. GPT 4.5 and DeepSeek R1 are still much better. (Maybe OpenAI will just add a writing tool call that calls on a dedicated, writing model - they have [teased their Creative Writing model](#) and we'd really like to see it!)

Looking at business writing like improving my LinkedIn posts, GPT-4.5 stays much truer to my tone and gives me portions of text I'd actually use, vs. GPT-5's more 'LinkedIn-slop' style response.



Here's an improved and more engaging version of your LinkedIn post:

GPT 4.5

It's been amazing to see all the excitement around **PII Guard**!

At Raindrop, we help you detect issues in AI agents—without compromising user privacy. **PII Guard** lets you accurately track problem while customizing exactly what personally identifiable information you redact, such as names, locations, medical or mental health details, proprietary code, API keys, and more.

Accurate AI monitoring meets robust user privacy. No compromises required.



GPT 5

PII Guard is live in Raindrop.

- No trade-off between privacy and detection quality
- Granular, configurable redaction (names, locations, medical/mental health info, proprietary code, API keys, ...)
- Works across your agent evaluations and monitoring Want a quick walkthrough? I'll send the demo link—just reply "PII."

Optional hooks/taglines you can use at the top

- Spot issues. Not identities.
- Keep what matters. Redact what doesn't.
- Privacy without blind spots for AI agents.

Hashtags (pick 3–5)

#AI #Security #Privacy #PII #DataProtection #AIAgents #AIObservability

While I never use AI for personal writing (because I have a strong belief in writing that is truly human, I think), I was curious how 4.5 and 5 would stack up playing with less structured content. While neither are great, 4.5 once again stays much truer to my tone and sounds far less like LLM slop than its counterparts.

ChatGPT 4.5 >



Help me improve this writing: AGI is an excellent vision of the future. One of the best of all time. It's generations of human imagination literally compressed into 3 letters. The countless movies, books, songs, essays, dreams about intelligence that meets and surpasses our own have given us a vision to turn into reality. The Turing Test and newer evals are even clearer articulations of the future - specific expectations to achieve.

Here's an improved, refined version that preserves your core message while enhancing clarity and impact:

GPT 4.5

AGI is one of humanity's most compelling visions for the future. It compresses generations of imagination into three simple letters, encompassing countless movies, books, songs, essays, and dreams about intelligence meeting—and surpassing—our own. Benchmarks like the Turing Test and modern evaluations sharpen this vision, providing clear expectations and concrete milestones as we turn imagination into reality.




GPT 5

- AGI is a compelling vision of the future—three letters that condense generations of human imagination. Countless films, books, songs, essays, and dreams about intelligence that meet and surpasses—our own have given us a blueprint to pursue. The Turing Test and its modern successors translate that vision into measurable targets.

Final Thoughts

Our lived experience cosigns the official benchmarks that OpenAI published today



I think GPT-5 is unequivocally the best coding model in the world. We were probably around 65% of the way through automating software engineering, and now we *might* be around 72%. To me, it's the biggest leap since 3.5 Sonnet.

I'm extremely curious to find out how everyone else will receive the model. My guess is that most non-developers will not get it for a few months. We'll have to wait for the models to be integrated into products.

What's next?

Well... Sam's [2 year old todo list](#) is still outstanding...



Subscribe to Latent.Space

Thousands of paid subscribers

The AI Engineer newsletter + Top 10 US Tech podcast. Exploring AI UX, Agents, Devtools, Infra, Open Source Models. See <https://latent.space/about> for highlights from Chris Lattner, Andrej Karpathy, George Hotz, Willison, Soumith Chintala et al!

Subscribe

By subscribing, I agree to Substack's [Terms of Use](#), and acknowledge its [Information Collection Notice](#) and [Privacy Policy](#).



161 Likes · 24 Restacks

← Previous

Next



A guest post by

alexis

Subscribe to alexis



A guest post by

Ben Hylak

cto @ raindrop.ai ex-apple design, spacex avionics

Subscribe to Ben Hylak

Discussion about this post

Comments

Restacks



Write a comment...



Sahar Mor  Aug 7

 Liked by alexis, Latent.Space

It'll be interesting to see how OpenAI's web-first Codex experience will compare to Claude Cod first approach. Codex leans into onboarding and accessibility, while Claude prioritizes precision local context.

Both will push the industry into the direction of "build your own software".

 LIKE (8)  REPLY



TheNeverEndingFall Aug 8

 Liked by Latent.Space

Is it even worth learning to become a software engineer at this point?

 LIKE (4)  REPLY

6 replies

32 more comments...

© 2025 Latent.Space · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great culture